

ECEEC-355: Computer Organization & Architecture

Dr. Anup K. Das

*Electrical and Computer Engineering
Drexel University*

Content


- **Based on Chapter 4 of P&H**

Designing the microarchitecture

- **Blackbox**
- **Graybox**
- **Whitebox**

ISA to Microarchitecture (the blackbox)

Observe the direction of the arrows




Memory

Microarchitecture (the blackbox)

Lets take one such instruction


00110001100010101111110000101111



How many bits?

Microarchitecture (the blackbox)

Name (Field Size)	7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	Comments
R-type	funct7	rs2	rs1	funct3	rd	opcode	Arithmetic instruction format
I-type	immediate[11:0]		rs1	funct3	rd	opcode	Loads & immediate arithmetic
S-type	immed[11:5]	rs2	rs1	funct3	immed[4:0]	opcode	Stores
SB-type	immed[12,10:5]	rs2	rs1	funct3	immed[4:1,11]	opcode	Conditional branch format
UJ-type	immediate[20,10:1,11,19:12]				rd	opcode	Unconditional jump format
U-type	immediate[31:12]				rd	opcode	Upper immediate format



Microarchitecture (the blackbox)

- **What is opcode?**
 - It tells you what is going to be the operation

RISC-V-Opcodes

Operations	Opcode
add	0110011
addi	0010011
..	
...	
...	
...	
...	

What do the processor do?

Performs some operation on data

How to know which operation to perform?


From the opcode

Microarchitecture (the graybox)

- **Once we have retrieved an instruction, we need to decode the opcode**
- **Steps so far**
 - Fetch the instruction from instruction memory
 - Decode the opcode of the instruction
 - What is the next step?
 - Perform the desired operation
- **Microarchitectural steps for instruction processing**
 - Fetch → Decode → Execute → ...

Microarchitecture (the graybox)

- **Before we move on**
 - Few quick notes on the instruction fetch



Microarchitecture (the graybox)

- **Recollect: instructions need to be fetched sequentially (one after another)**
 - Control flow / Von-Neuman model
 - Unless you encounter conditional instructions
 - After executing the conditional instruction, you need to go back to the point where you were before to continue executing the next instruction
- **You need a register to keep track of the instruction you are executing**
 - Program counter register
 - This is the **first** component of your microarchitecture

Microarchitecture (the graybox)

- Recap of memory organization


a	a	a+1	a+2	a+3	a+4	a+5	a+6	a+7
a+8	a+8	a+9	a+10	a+11	a+12	a+13	a+14	a+15
a+16	a+16	a+17	a+18	a+19	a+20	a+21	a+22	a+23
a+24	a+24	a+25	a+26	a+27	a+28			
⋮								

- Remember each row is a double word (64 bits)**
 - Each instruction is 32 bits (occupies 4 blocks in a double word)
 - Each new instruction is offset by 4 (a, a+4, a+8, ...)

Microarchitecture (the graybox)

- **Instructions**
 - Sequential
 - Control
- **Lets first focus on sequential instructions**
 - Instructions that are fetched and executed one after the other
 - Control instructions jumps to an specified location in the code
 - if-else, while, for, etc.
 - Conditional/unconditional jumps (jal,jar,beq,etc.)
- **For sequential instructions**
 - Next instruction address = current instruction address + 4
 - We have seen this in the last slide!

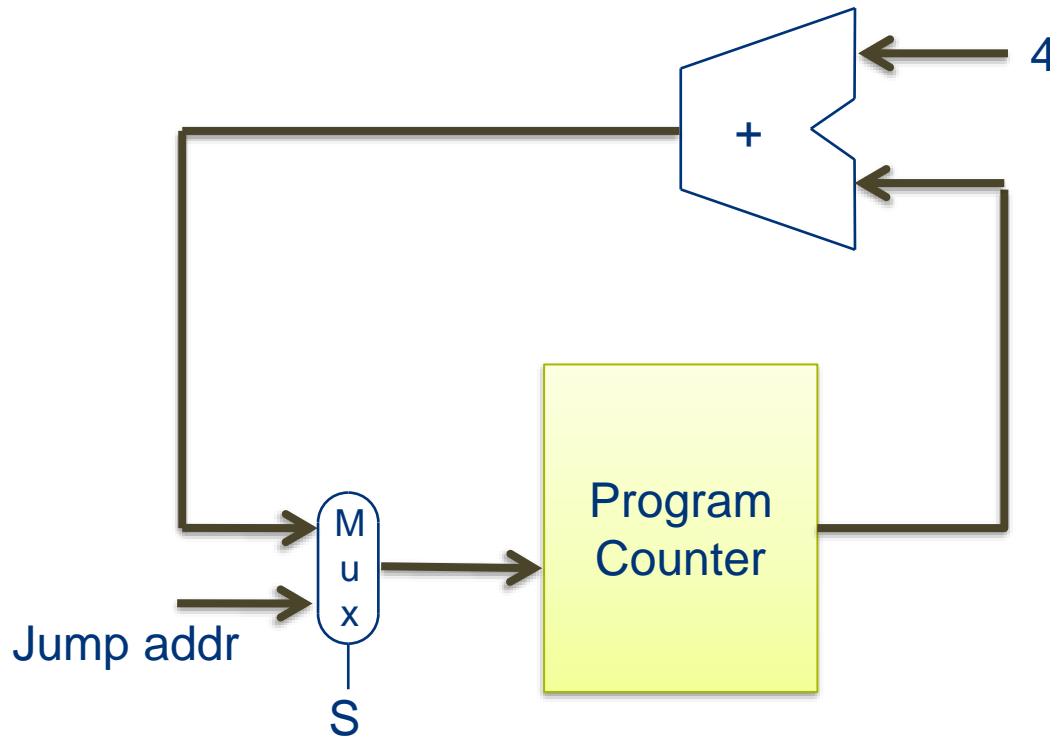
Microarchitecture (the graybox)



This is for the sequential part

We still need to have some logic for the conditional instructions

Microarchitecture (the graybox)

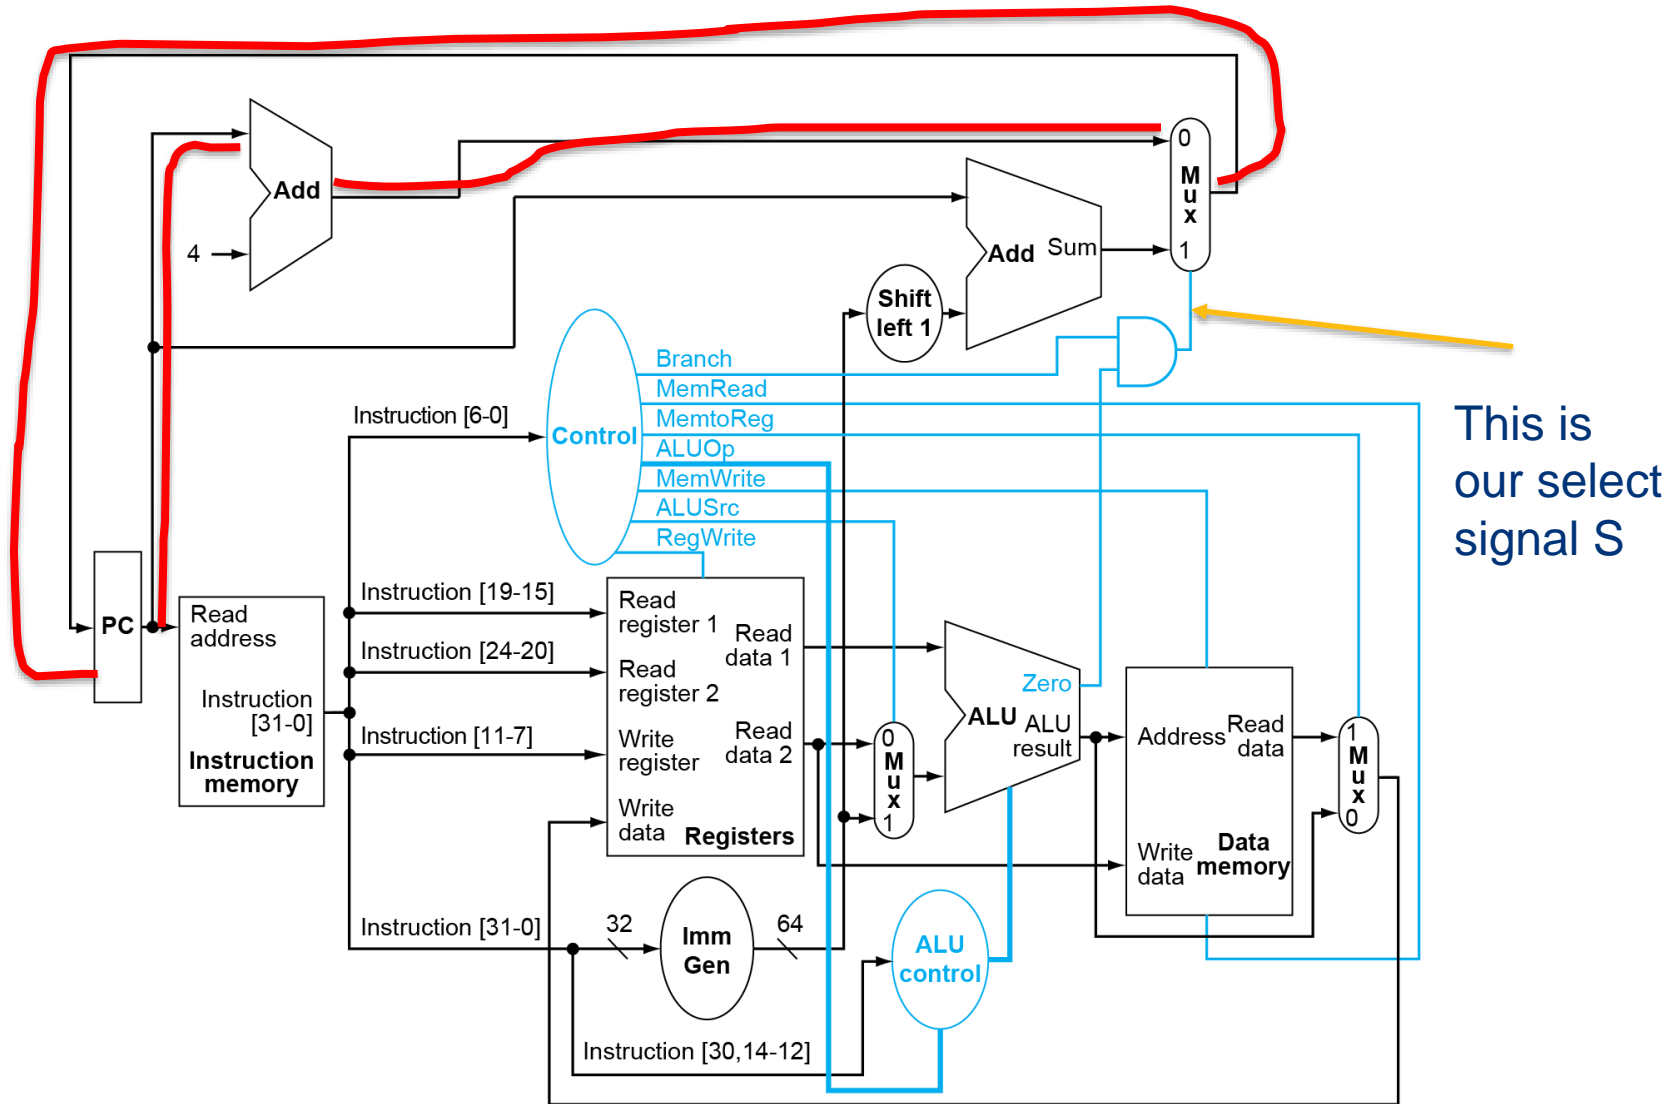


This is for the sequential part

We still need to have some logic for the conditional instructions

If $S = 0$: $pc = pc + 4$
If $S = 1$: $pc = \text{jump addr}$

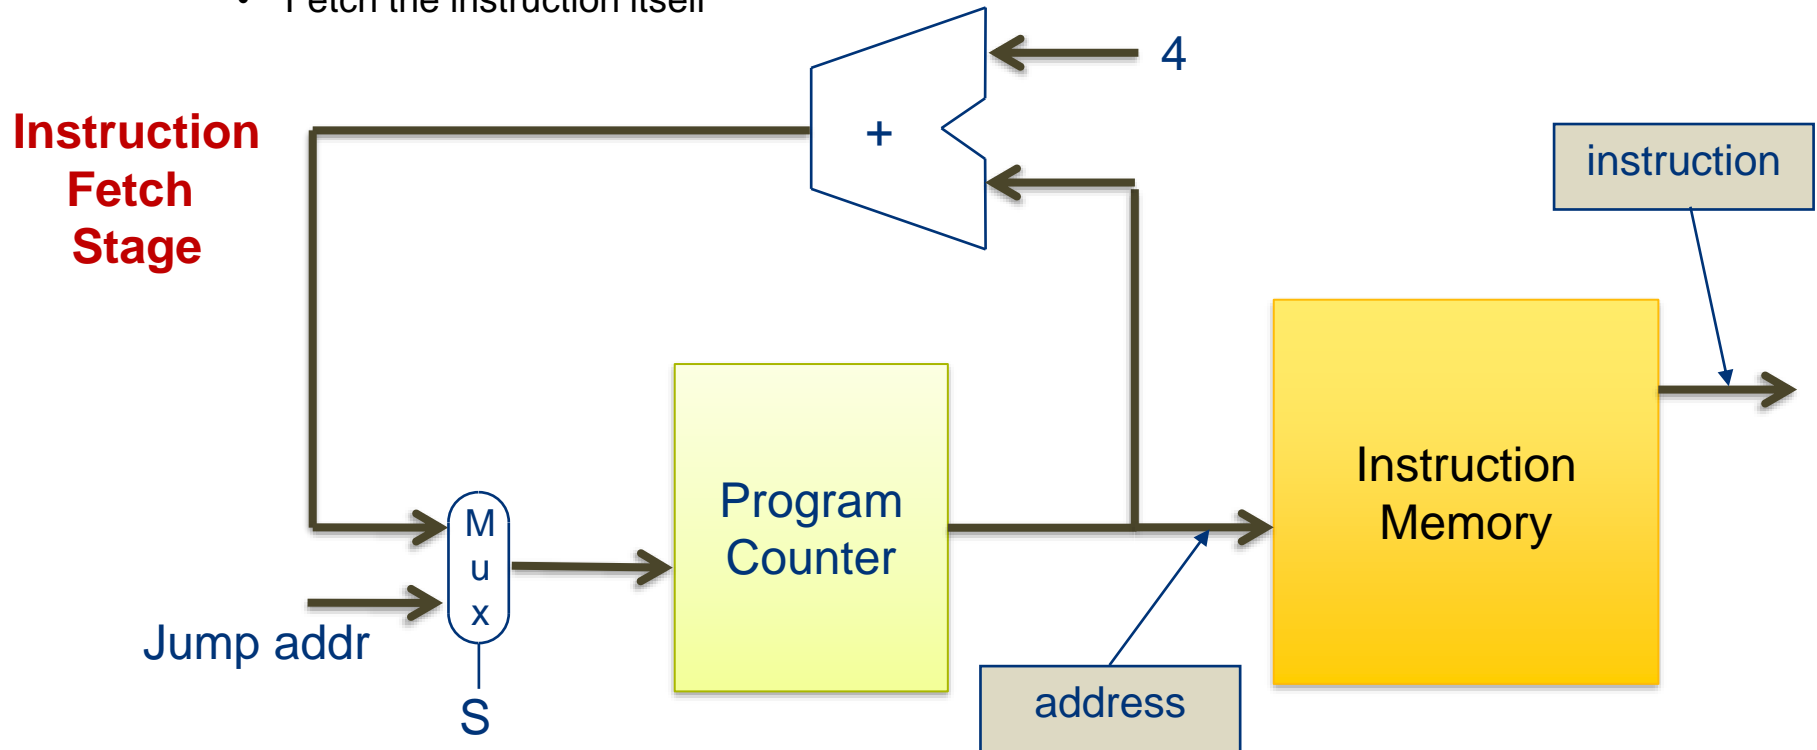
A quick look at the Microarchitecture



This is
our select
signal S

Microarchitecture (Gray to White)

- So we have the logic for fetching the **address** of the next instruction
 - What to do next?
 - Fetch the instruction itself



Microarchitecture (Gray to White)

- You have got the instruction, what to do next?
 - Decode
 - Extract certain fields from the instruction to
 - Find out what to do, i.e., operations
 - Where to do it to, i.e., operands

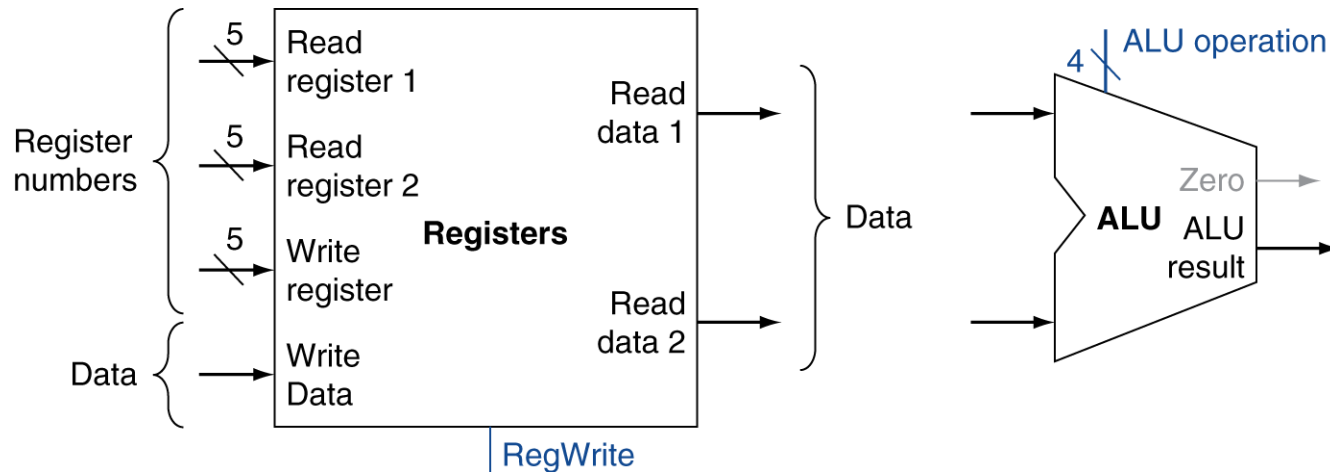
Name (Field Size)	Field						Comments
	7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	
R-type	funct7	rs2	rs1	funct3	rd	opcode	Arithmetic instruction format
I-type	immediate[11:0]		rs1	funct3	rd	opcode	Loads & immediate arithmetic
S-type	immed[11:5]	rs2	rs1	funct3	immed[4:0]	opcode	Stores
SB-type	immed[12,10:5]	rs2	rs1	funct3	immed[4:1,11]	opcode	Conditional branch format
UJ-type	immediate[20,10:1,11,19:12]				rd	opcode	Unconditional jump format
U-type	immediate[31:12]				rd	opcode	Upper immediate format

Microarchitecture (Gray to White)

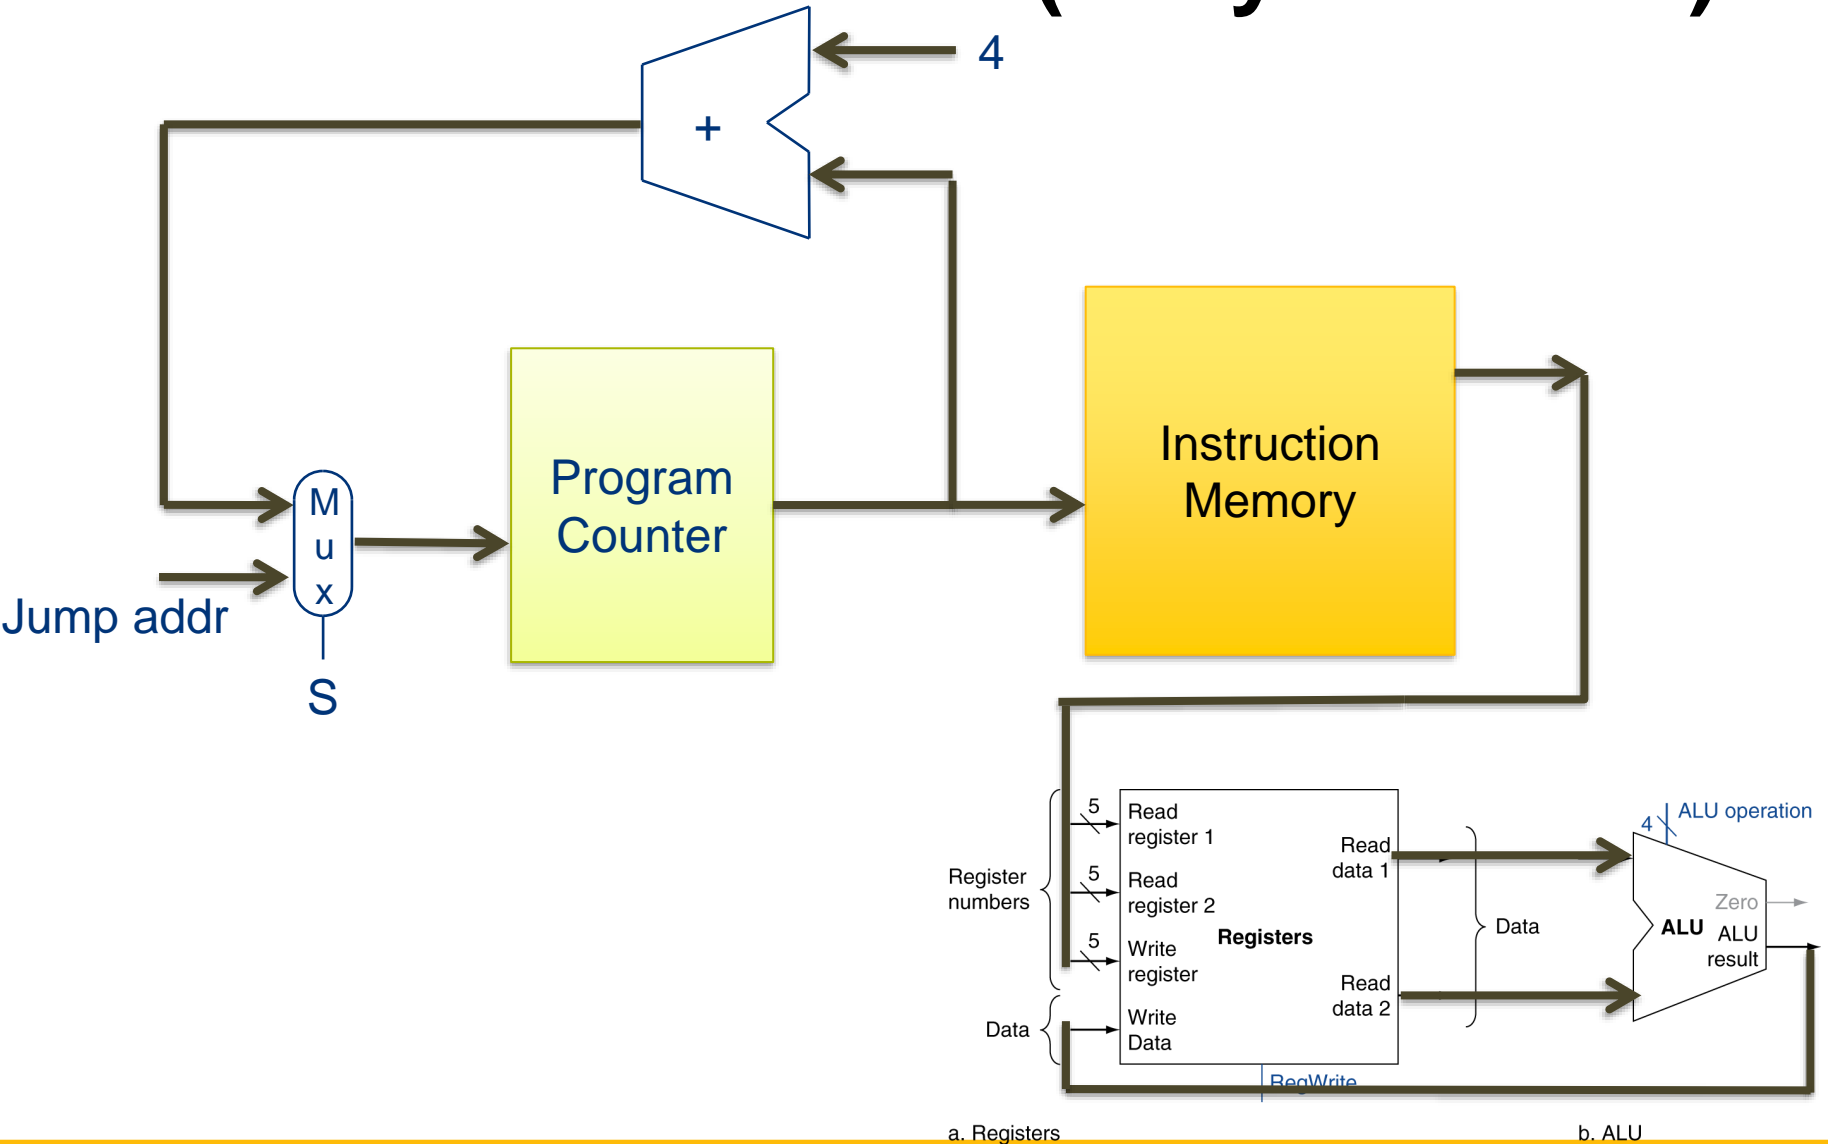
Name (Field Size)	7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	Comments
R-type	funct7	rs2	rs1	funct3	rd	opcode	Arithmetic instruction format
I-type	immediate[11:0]		rs1	funct3	rd	opcode	Loads & immediate arithmetic
S-type	immed[11:5]	rs2	rs1	funct3	immed[4:0]	opcode	Stores
SB-type	immed[12,10:5]	rs2	rs1	funct3	immed[4:1,11]	opcode	Conditional branch format
UJ-type	immediate[20,10:1,11,19:12]				rd	opcode	Unconditional jump format
U-type	immediate[31:12]				rd	opcode	Upper immediate format

• R-type Instructions

- Read two register operands
- Perform arithmetic/logical operation
- Write register result

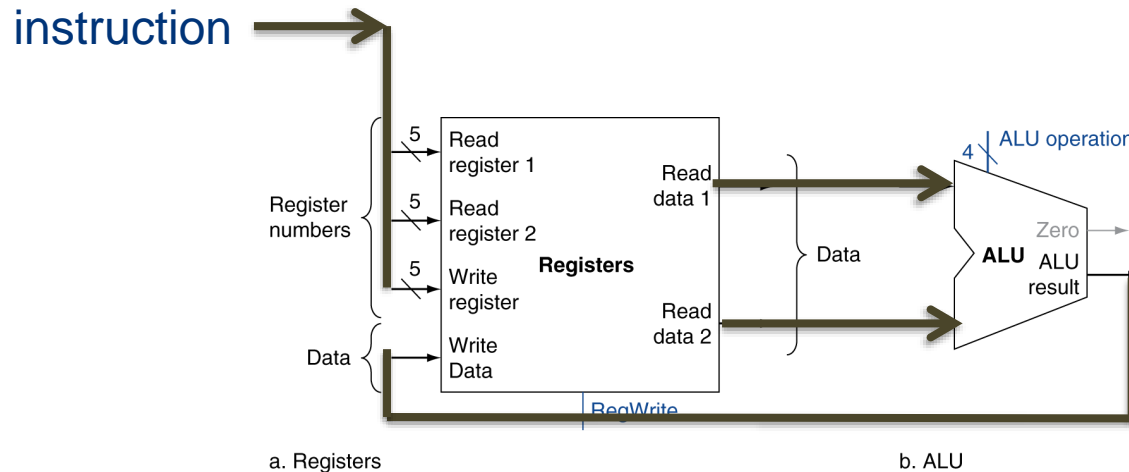


Microarchitecture (Gray to White)



Microarchitecture (Gray to White)

- **Every instruction must have fetch unit**
 - What differs is how the instruction is decoded
 - For simplicity we will skip the drawing the fetch unit




Microarchitecture (Gray to White)


Name (Field Size)	7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	Comments
R-type	funct7	rs2	rs1	funct3	rd	opcode	Arithmetic instruction format
I-type	immediate[11:0]		rs1	funct3	rd	opcode	Loads & immediate arithmetic
S-type	immed[11:5]	rs2	rs1	funct3	immed[4:0]	opcode	Stores
SB-type	immed[12,10:5]	rs2	rs1	funct3	immed[4:1,11]	opcode	Conditional branch format
UJ-type	immediate[20,10:1,11,19:12]				rd	opcode	Unconditional jump format
U-type	immediate[31:12]				rd	opcode	Upper immediate format

• Load/Store Instructions

- Read register operands
- Calculate address using 12-bit offset
 - Use ALU, but sign-extend offset
- Load: Read memory and update register
- Store: Write register value to memory

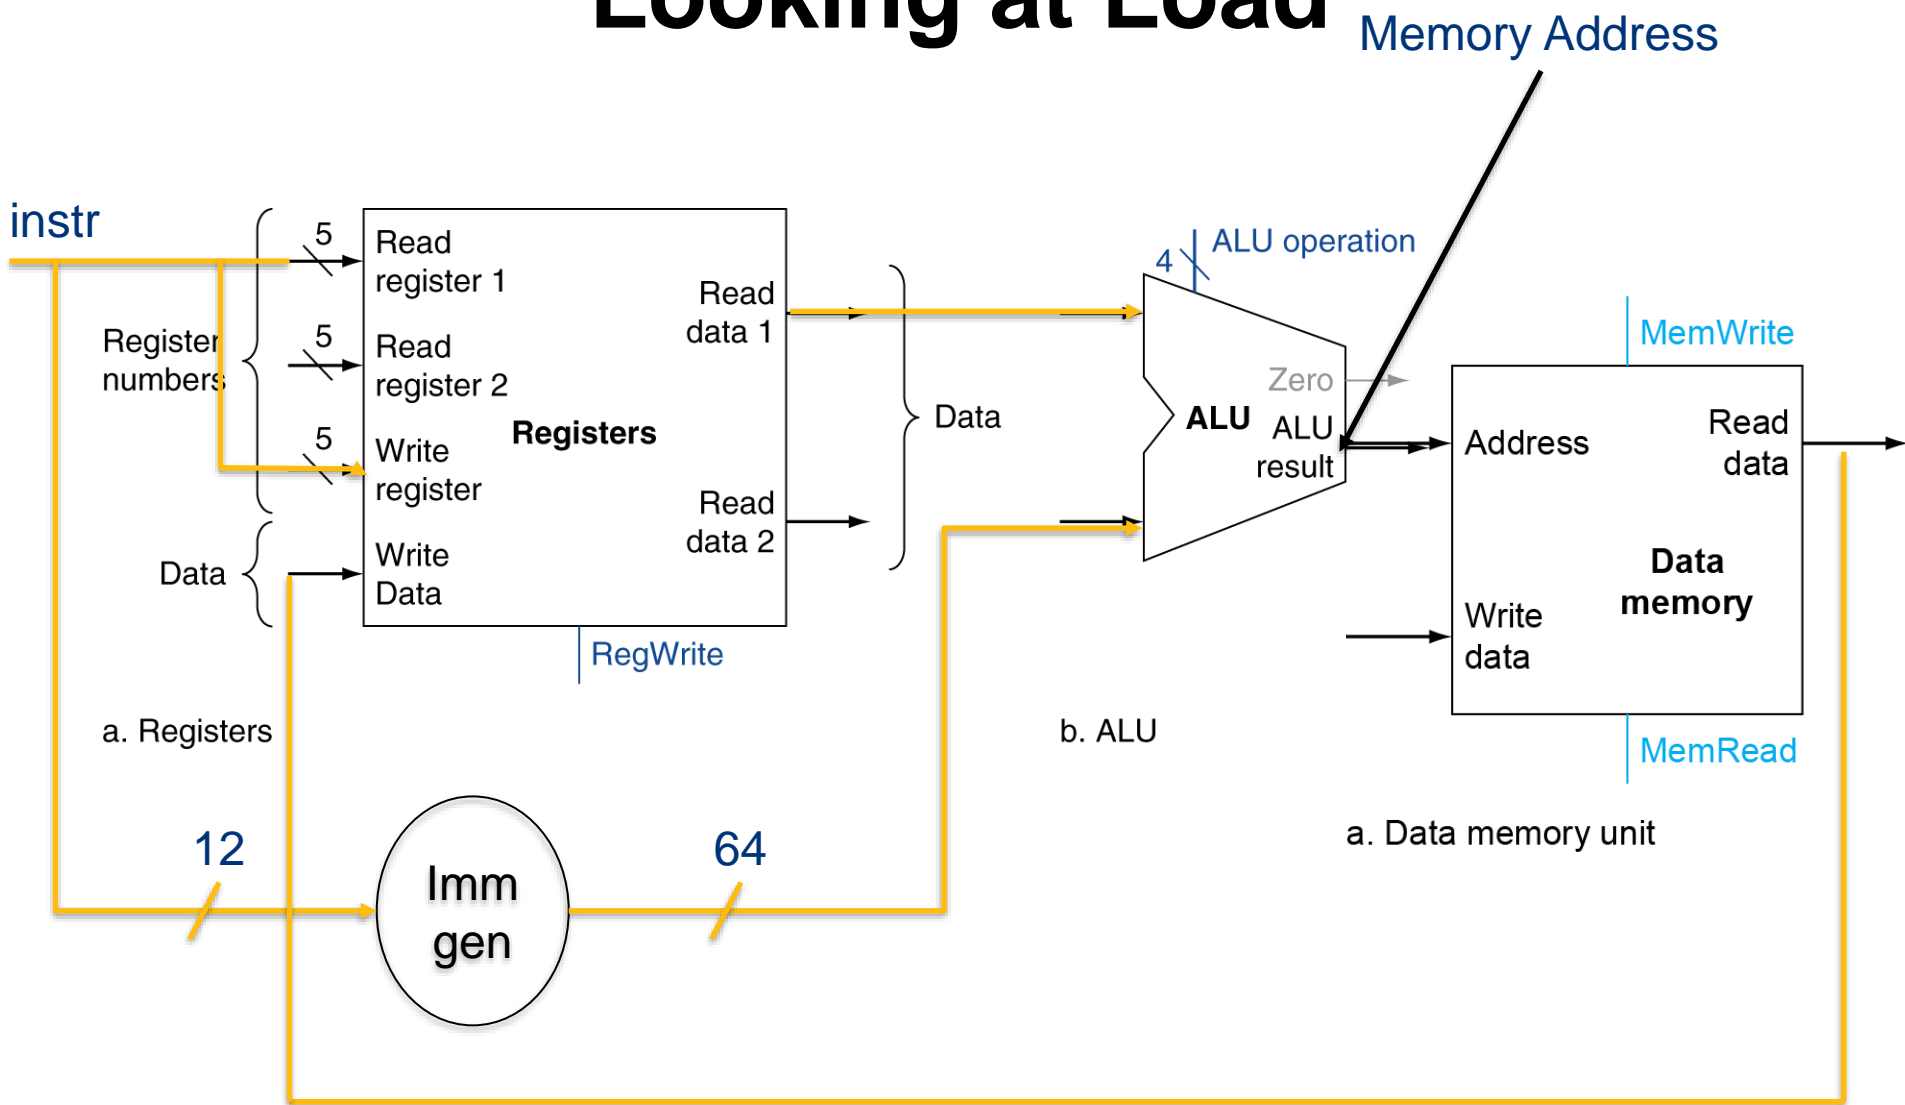


a. Data memory unit




b. Immediate generation

Looking at Load



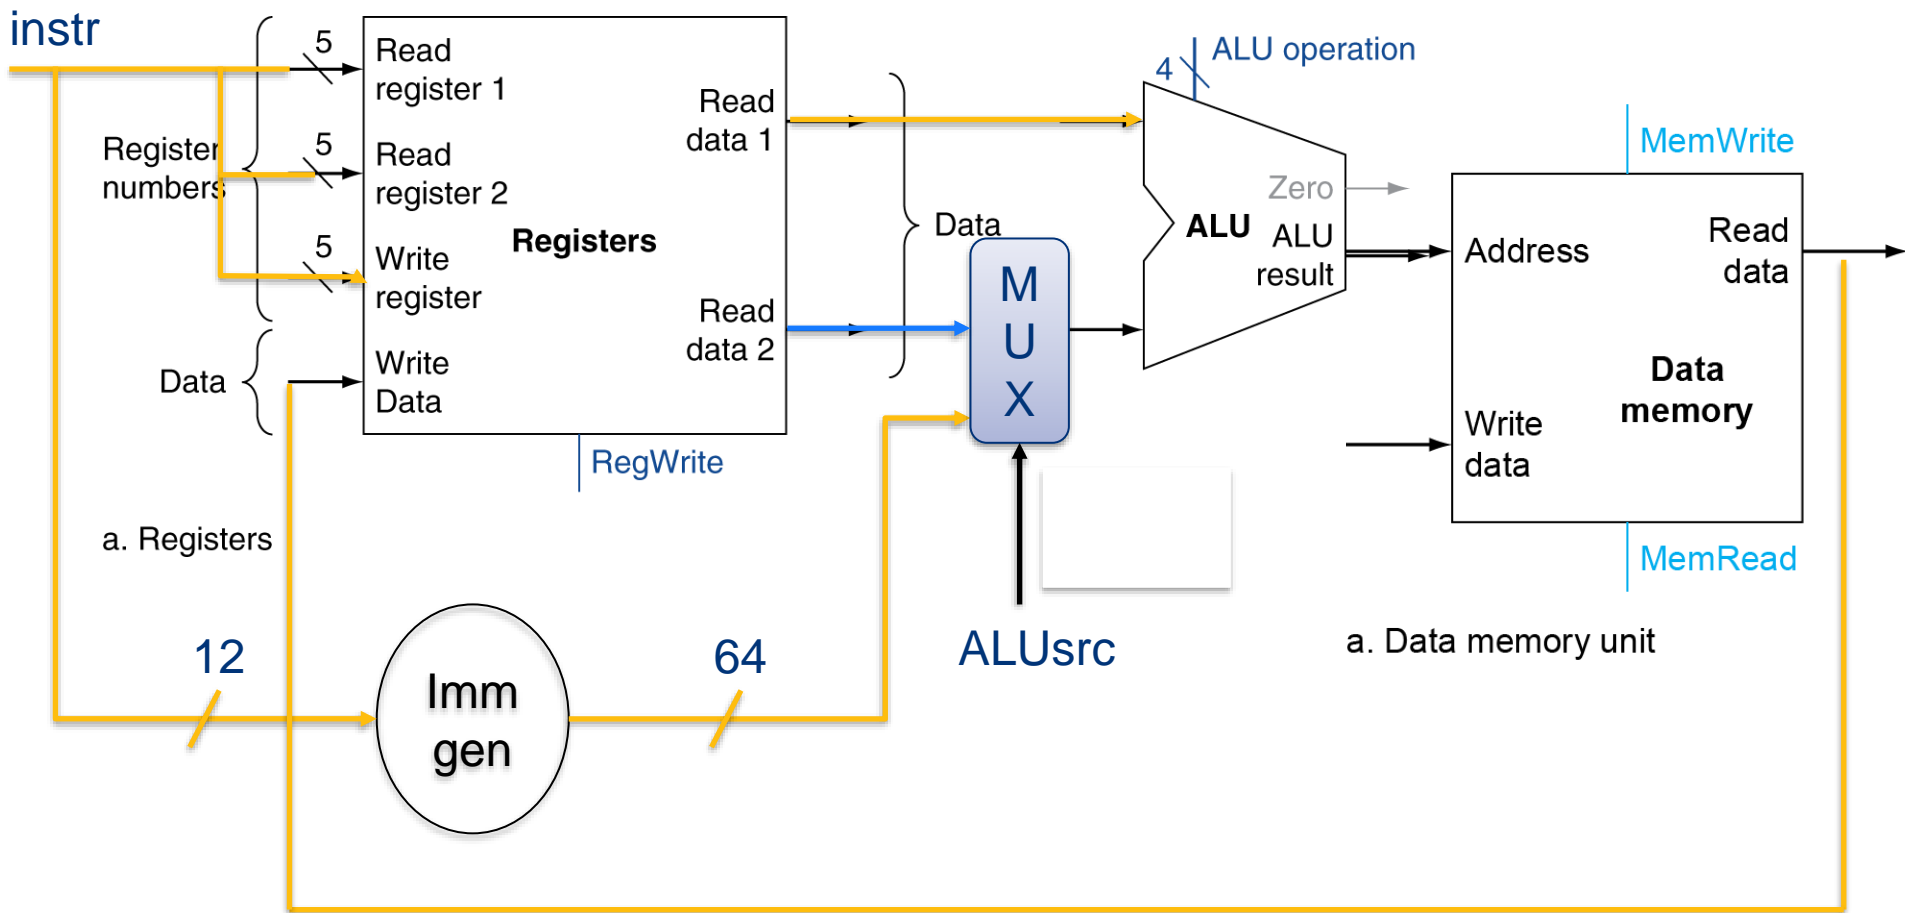
Observe the similarity with R-type



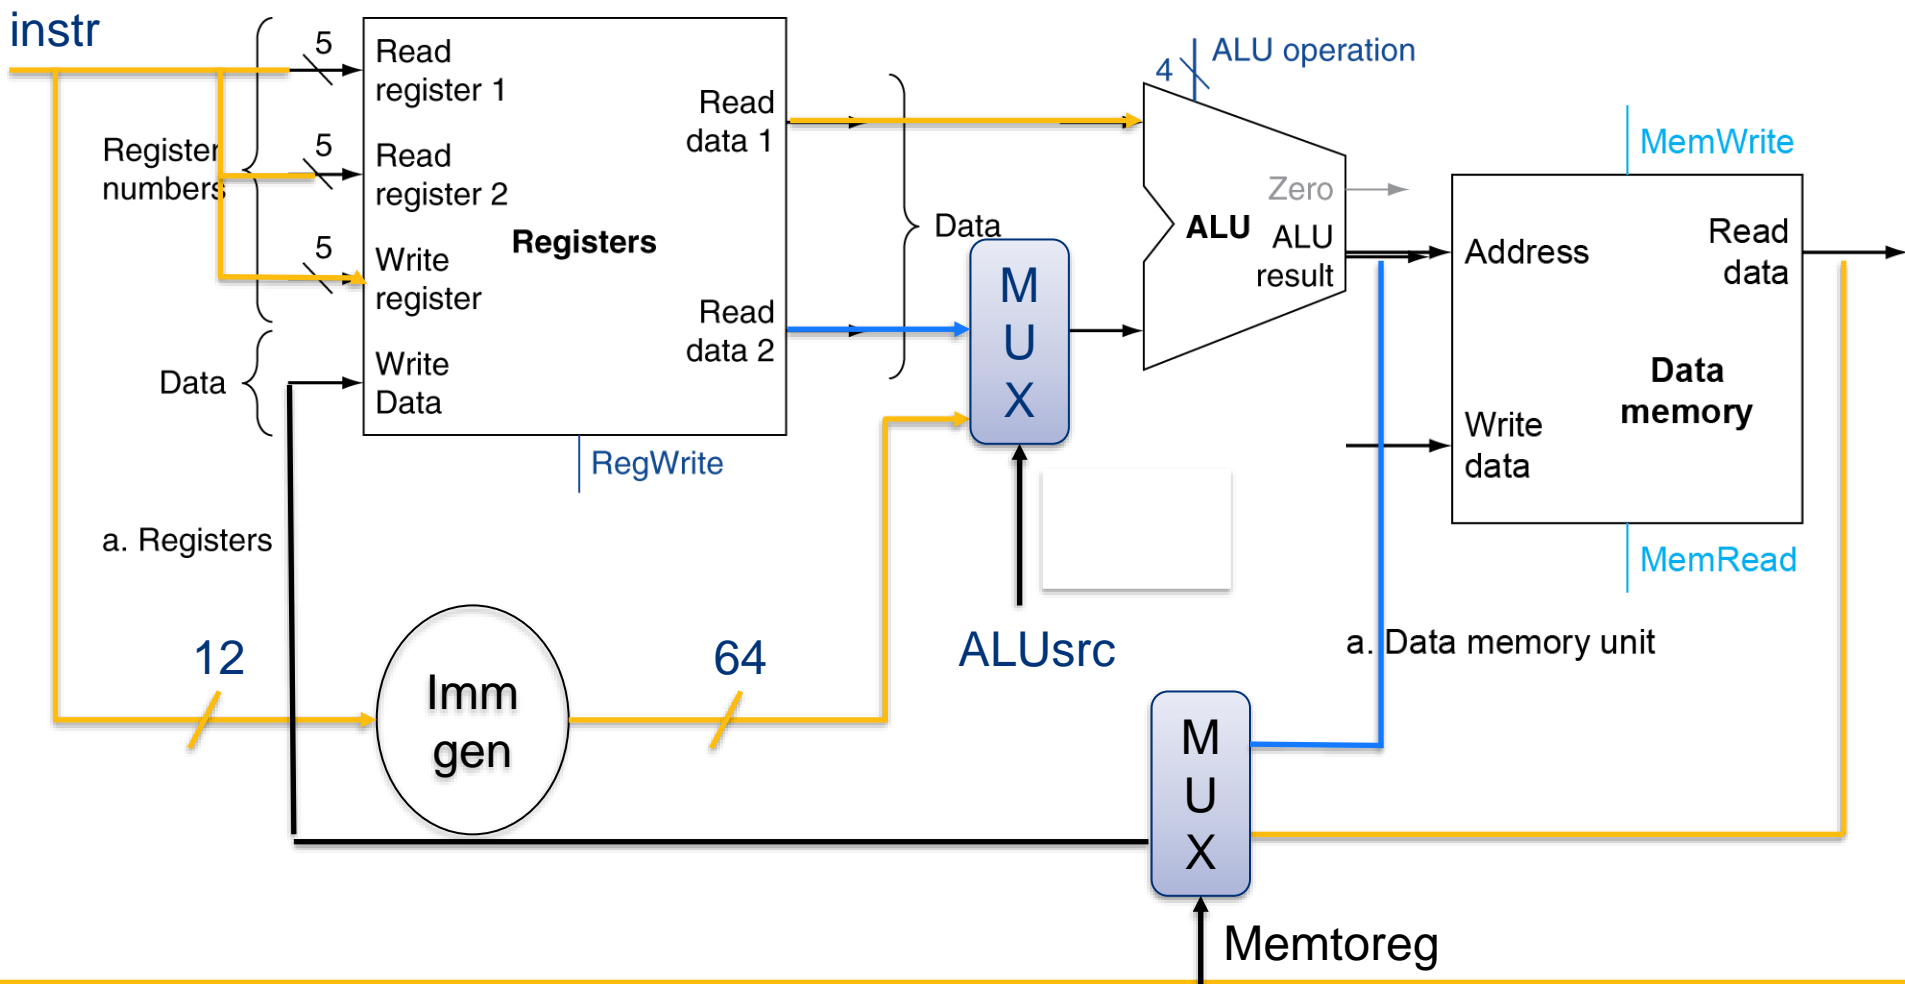
Observe Read data 2 goes to the second input of the ALU

Output of ALU goes back to the register

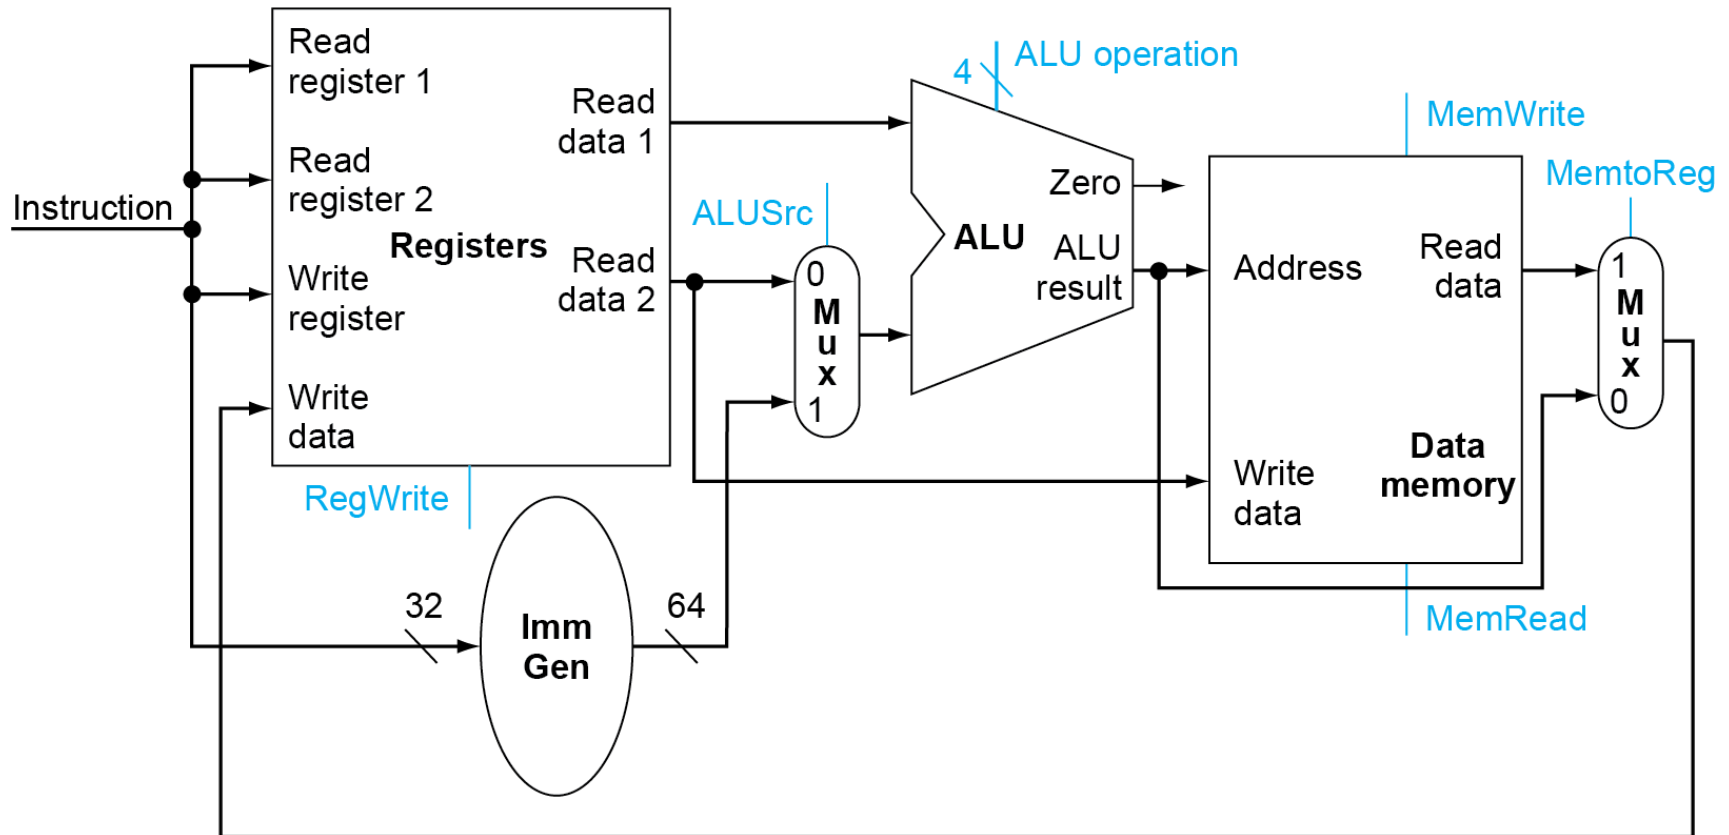
Looking at Load



Looking at Load



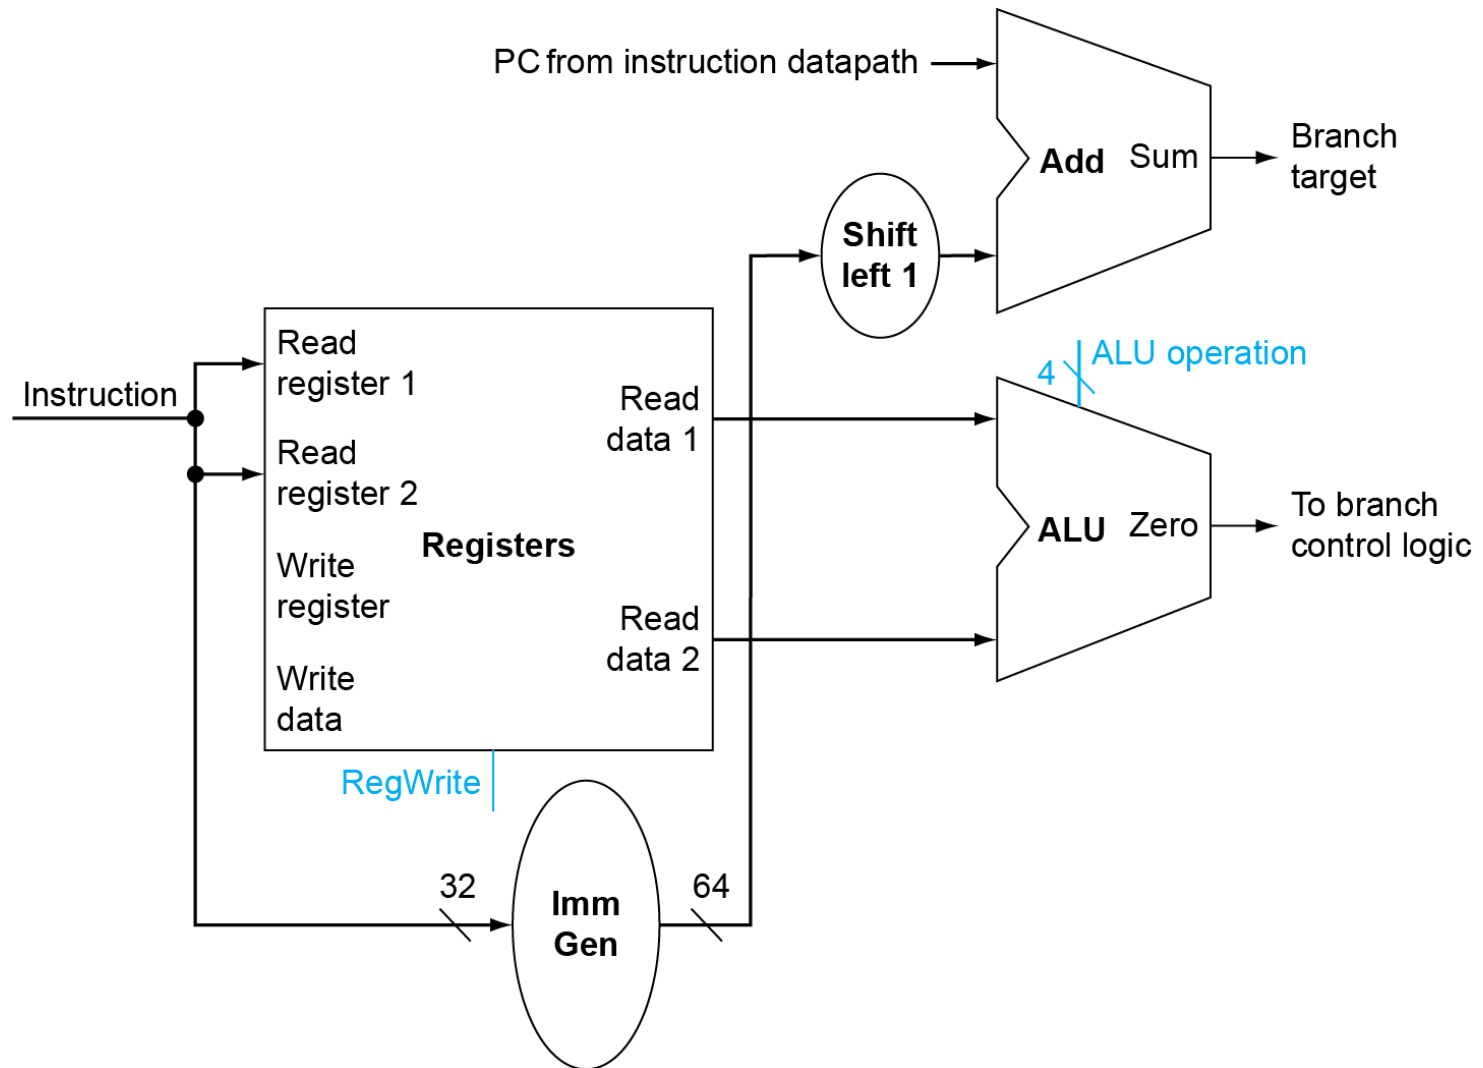
R-Type/Load/Store Datapath



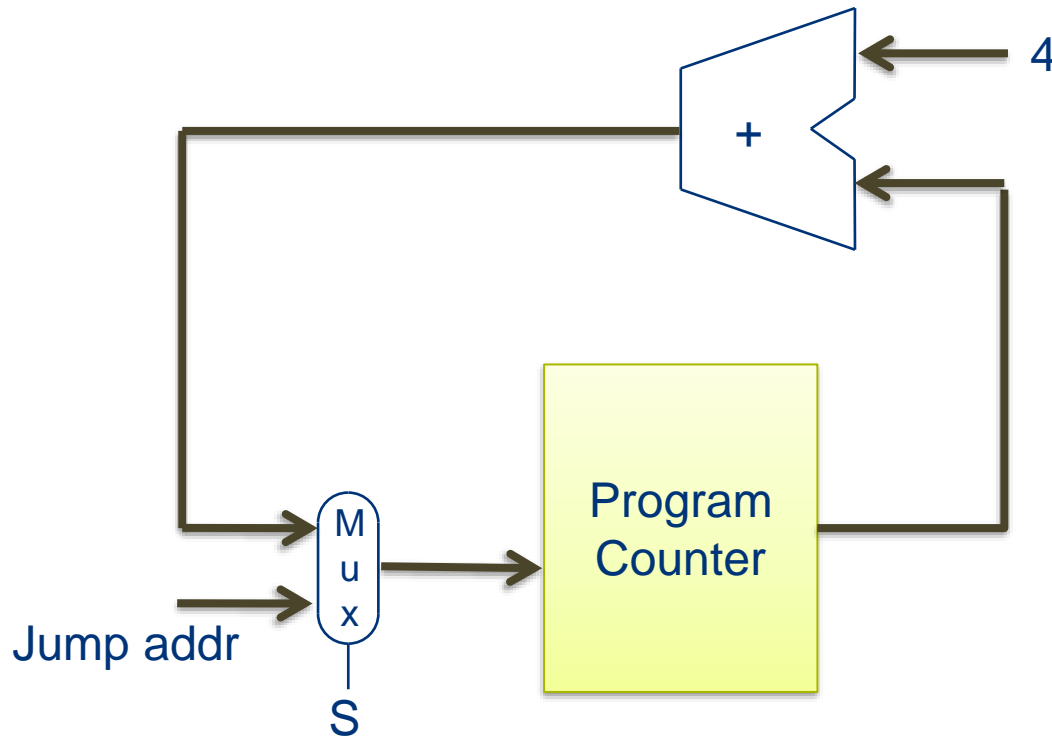
Branch Instructions

- **Read register operands**
- **Compare operands**
 - Use ALU, subtract and check Zero output
- **Calculate target address**
 - Sign-extend displacement
 - Shift left 1 place (halfword displacement)
 - Add to PC value

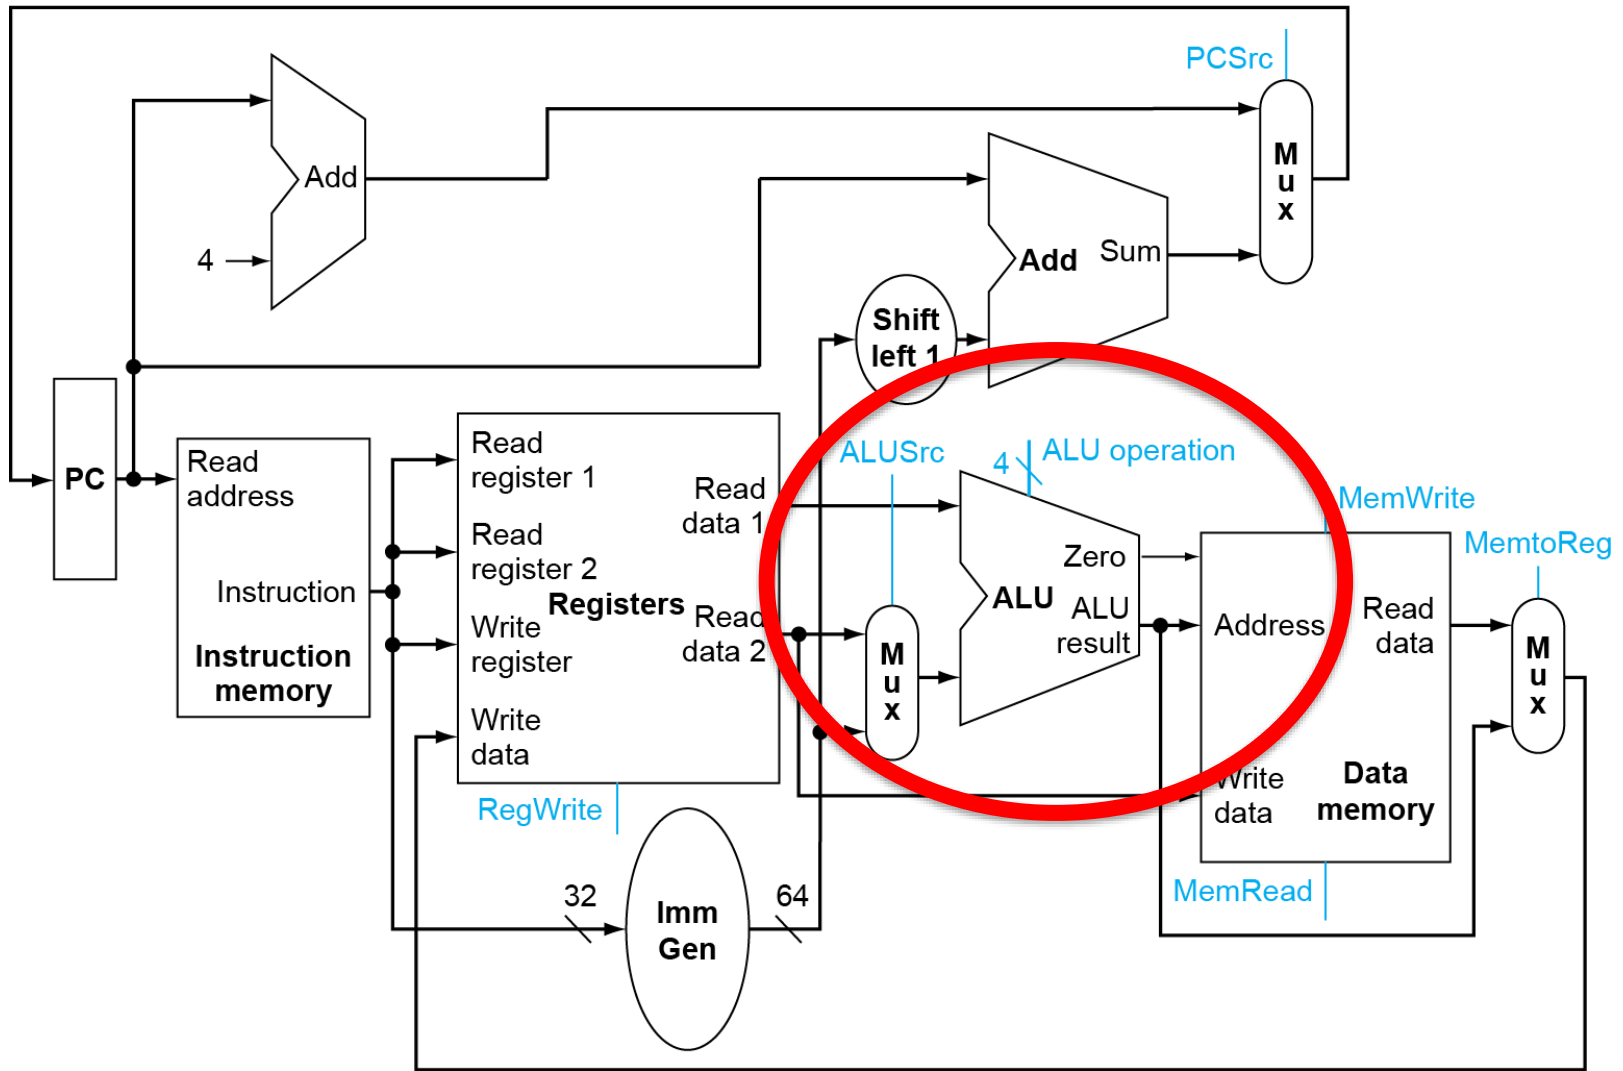
Branch Instructions



Remember PC?



Full Datapath



ALU Control

- **We build the input datapath to the ALU**
 - Basically using multiplexers
- **ALU themselves can perform lot more than addition**
 - This is controlled using 4 bits “ALU Operation”

ALU Control

- **ALU used for**
 - Load/Store: F = add
 - Branch: F = subtract
 - R-type: F depends on opcode

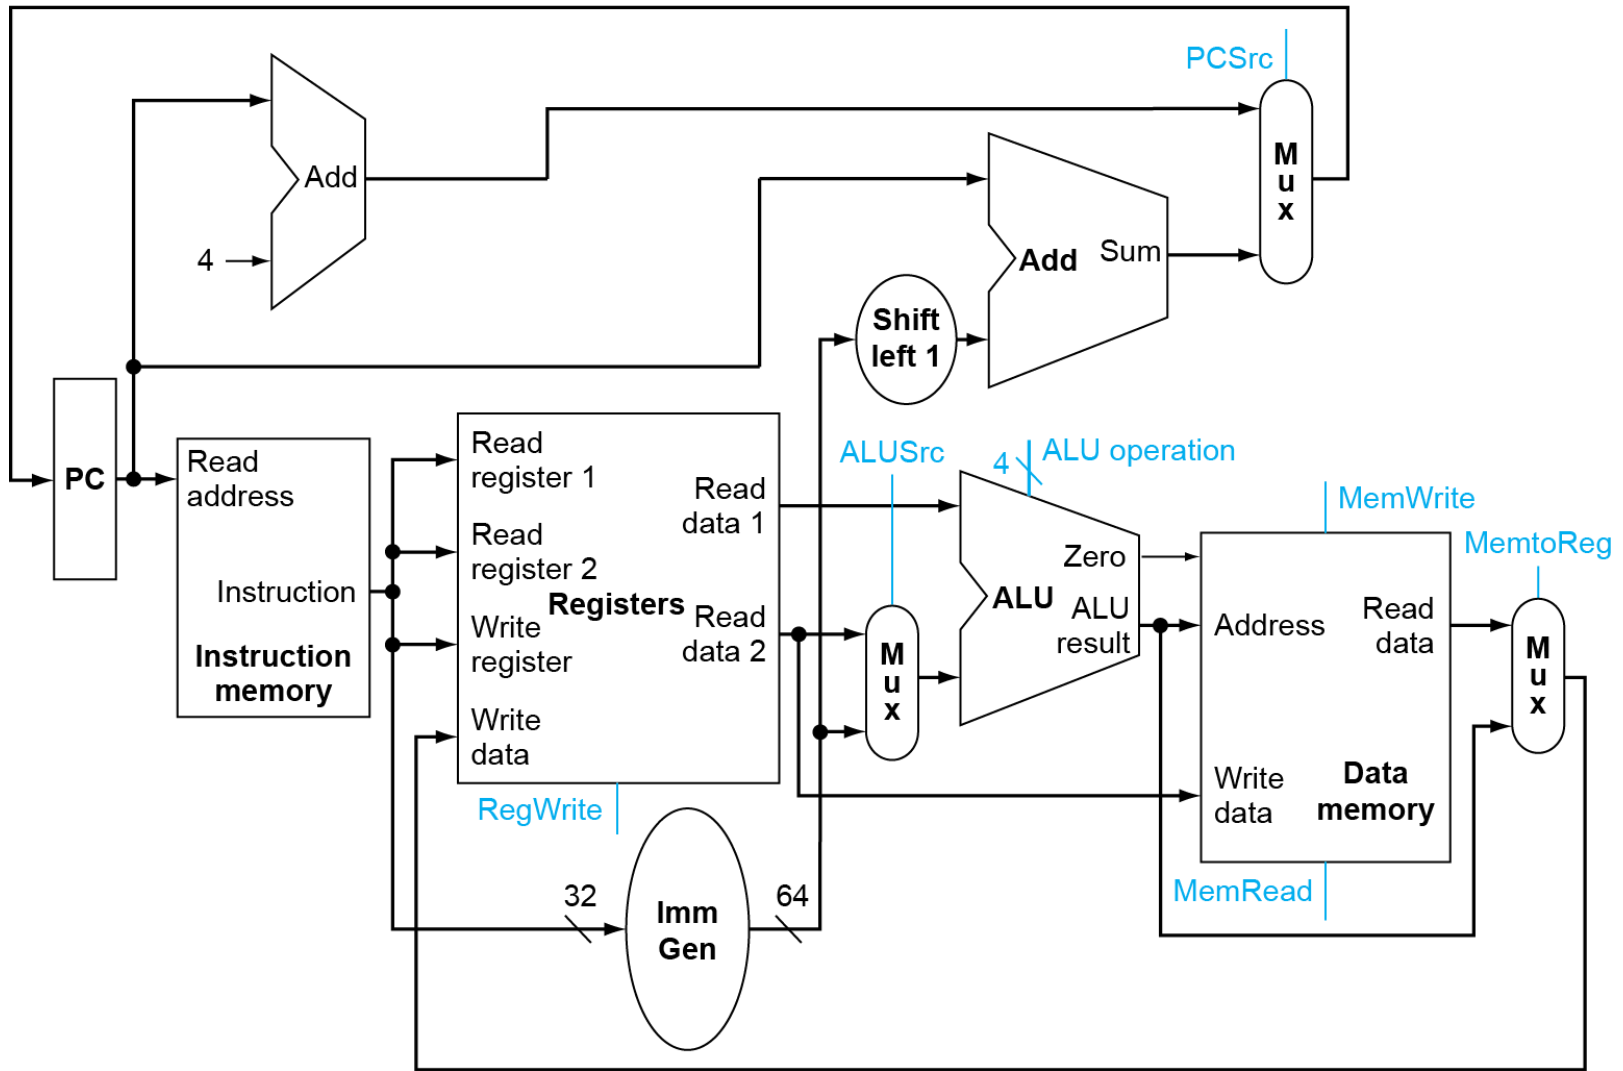
ALU control	Function
0000	AND
0001	OR
0010	add
0110	subtract

ALU Control

- **Assume 2-bit ALUOp derived from opcode**
 - Combinational logic derives ALU control

opcode	ALUOp	Operation	Opcode field	ALU function	ALU control
ld	00	load register	XXXXXXXXXXXX	add	0010
sd	00	store register	XXXXXXXXXXXX	add	0010
beq	01	branch on equal	XXXXXXXXXXXX	subtract	0110
R-type	10	add	100000	add	0010
		subtract	100010	subtract	0110
		AND	100100	AND	0000
		OR	100101	OR	0001

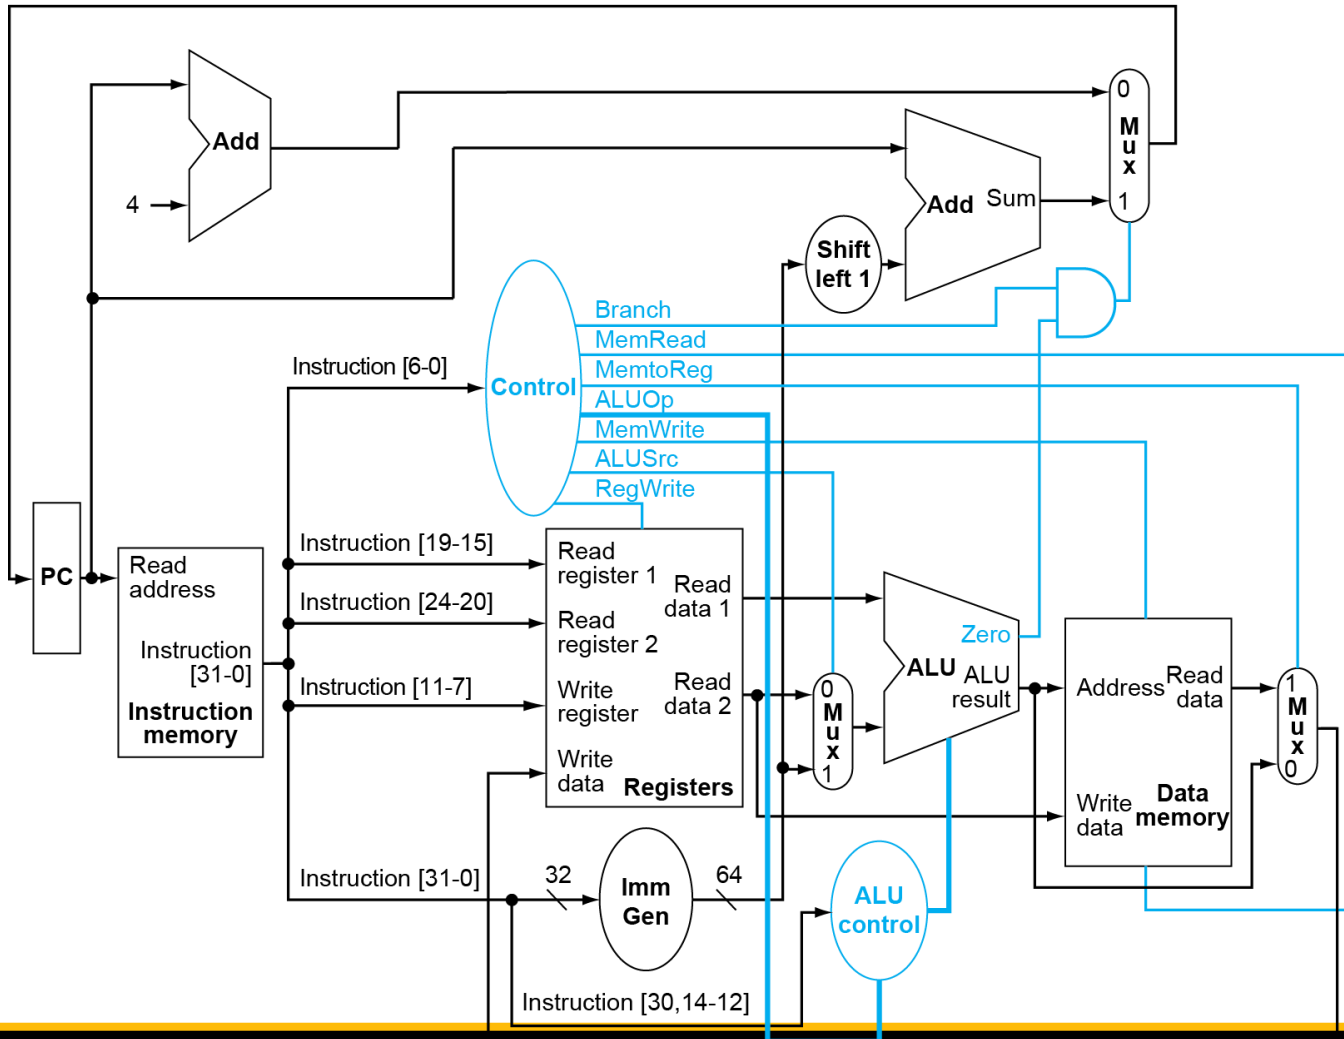
The Main Control Unit



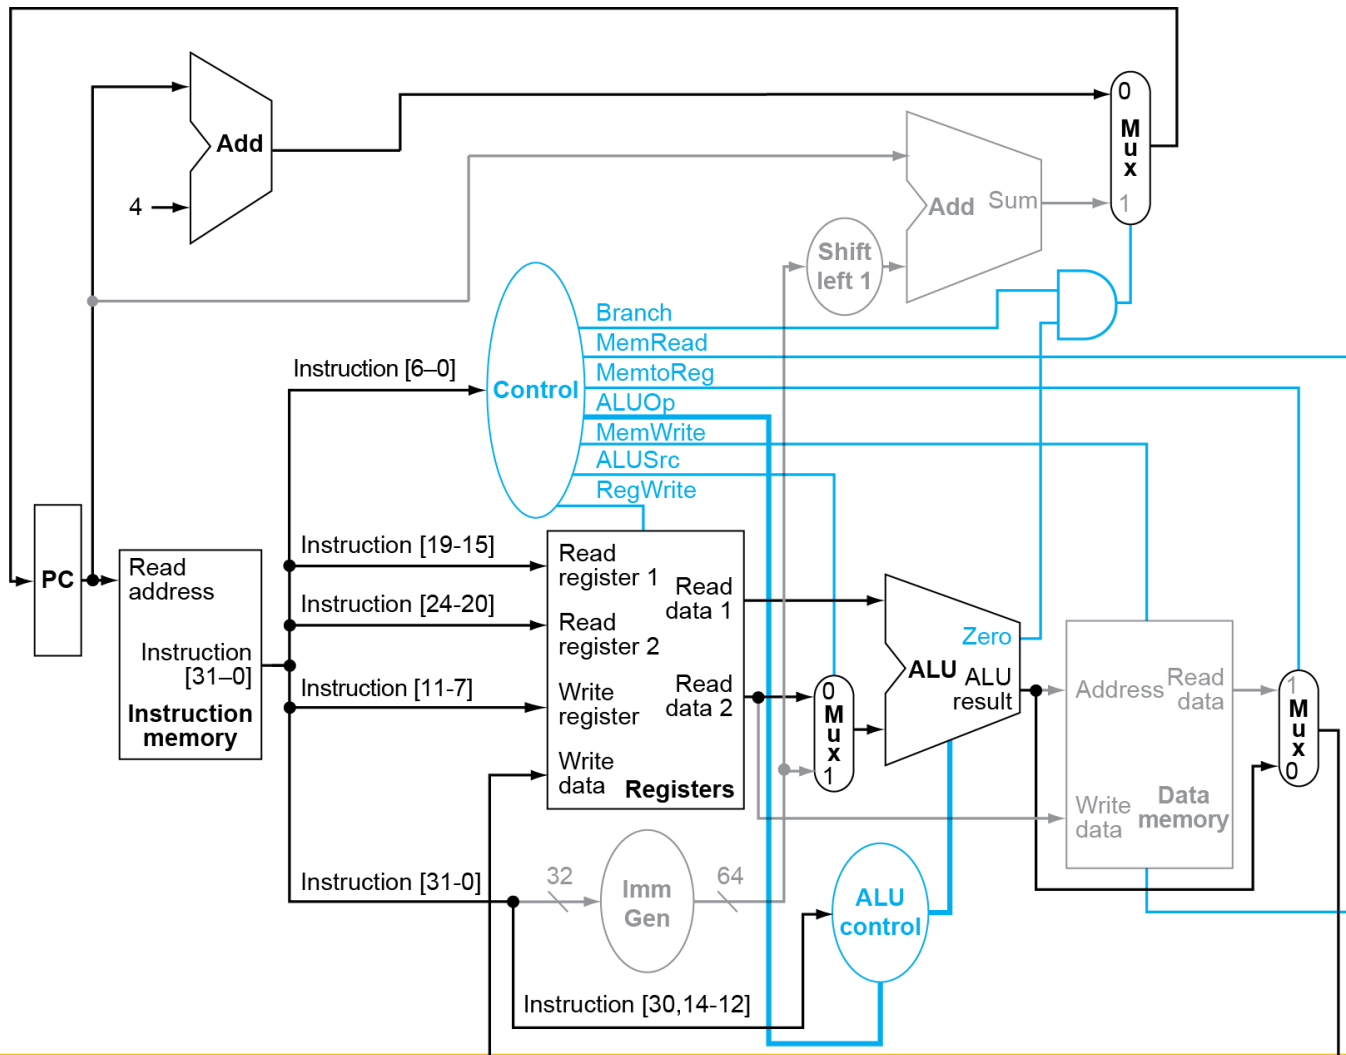
The Main Control Unit

- **They control the Multiplexers and functional units**
 - Based on the opcode

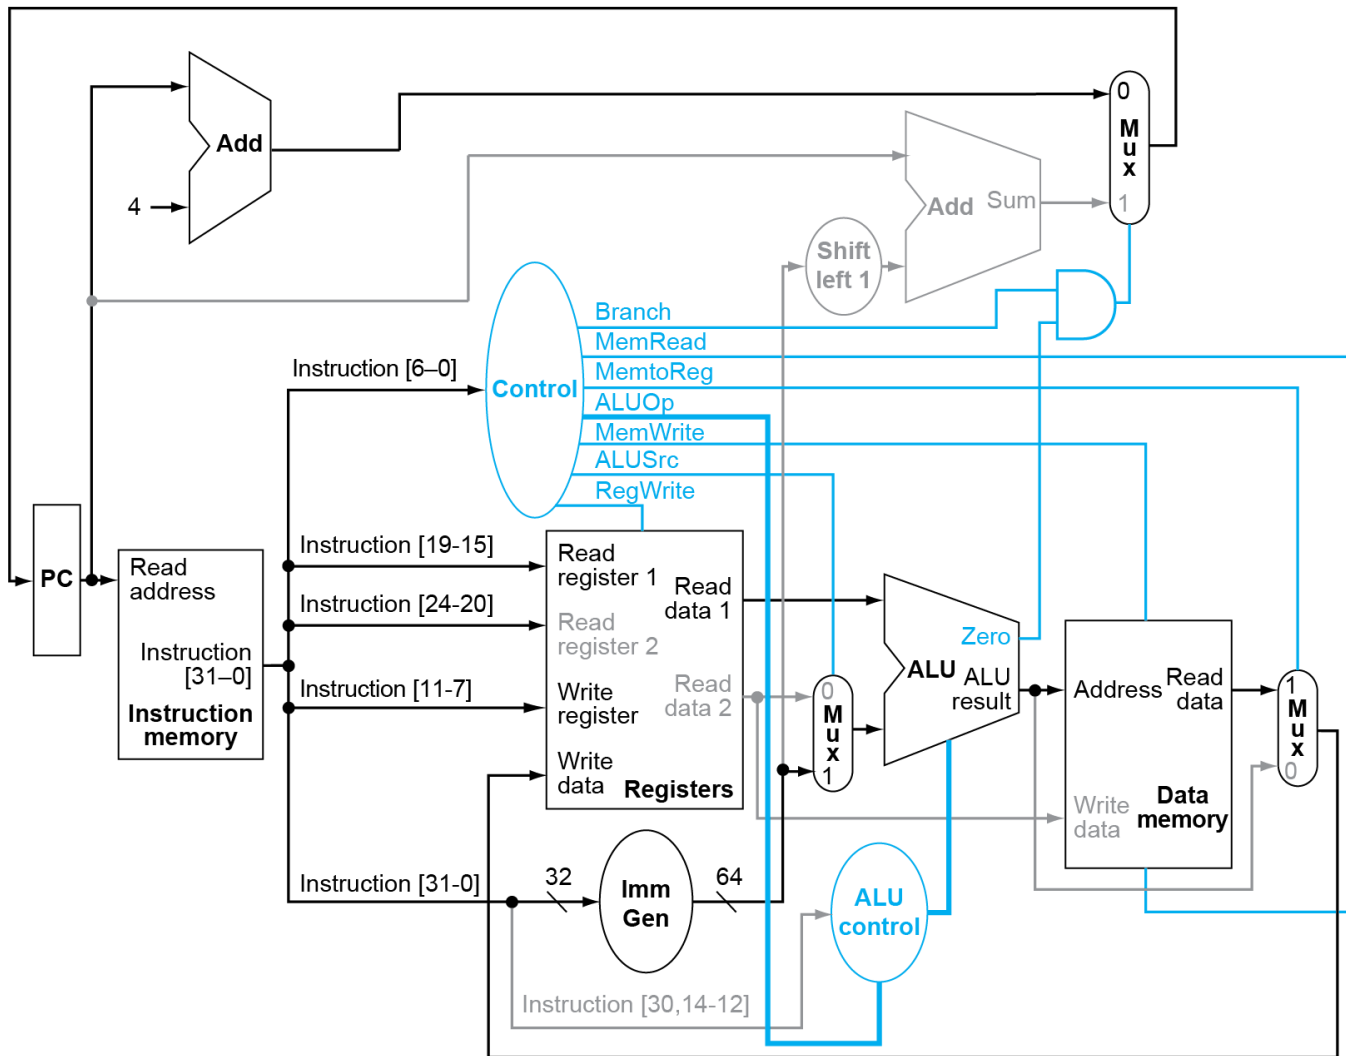
Datapath with control



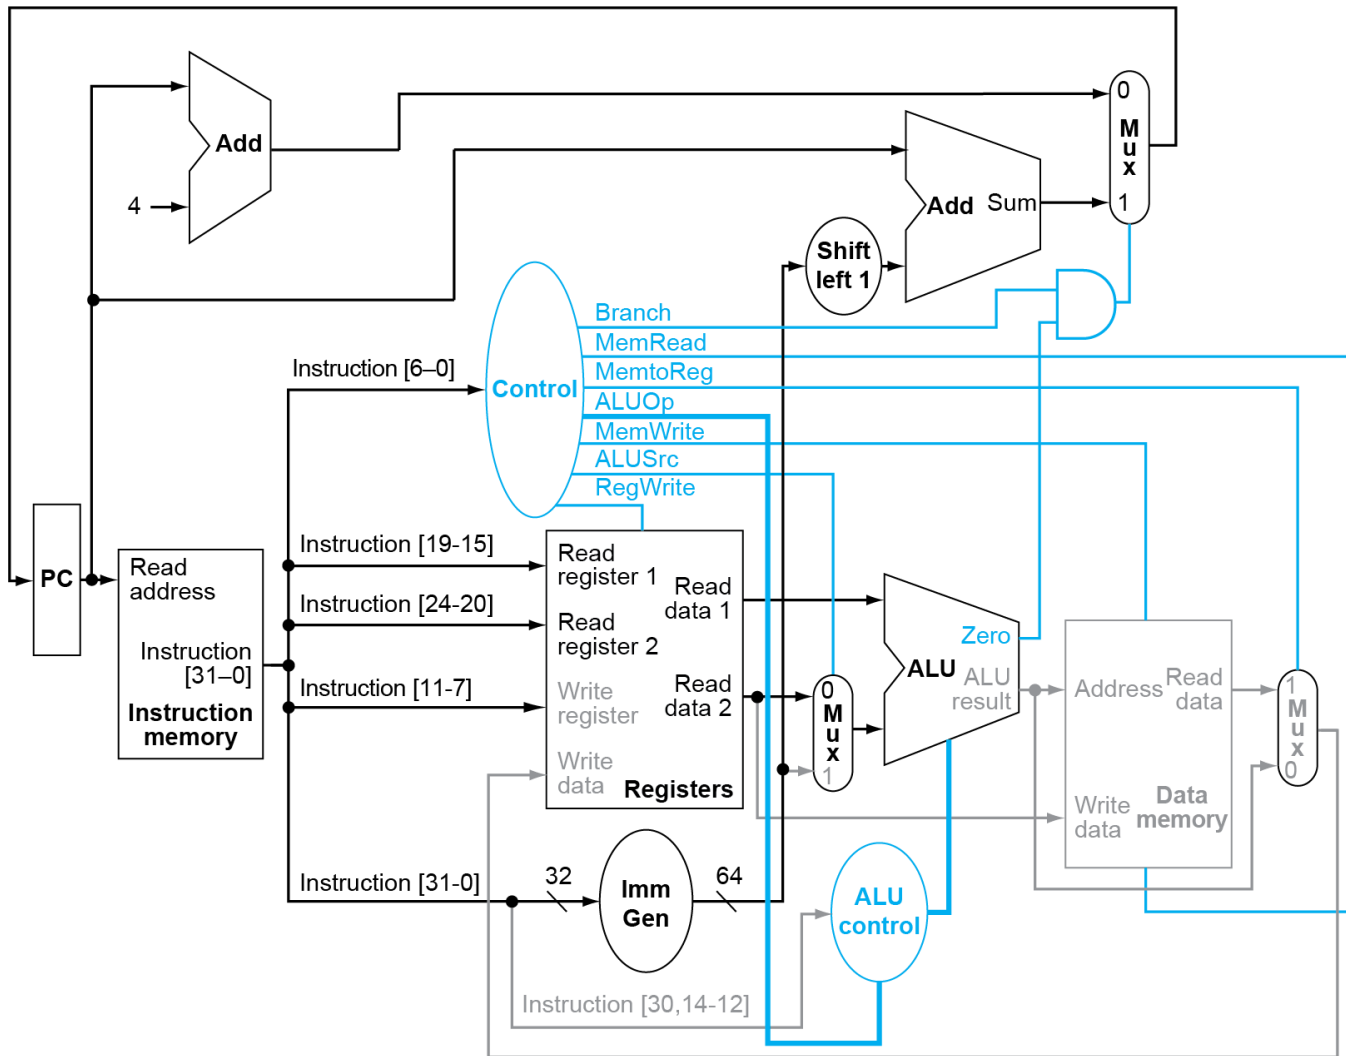
R-Type Instruction



Load Instruction



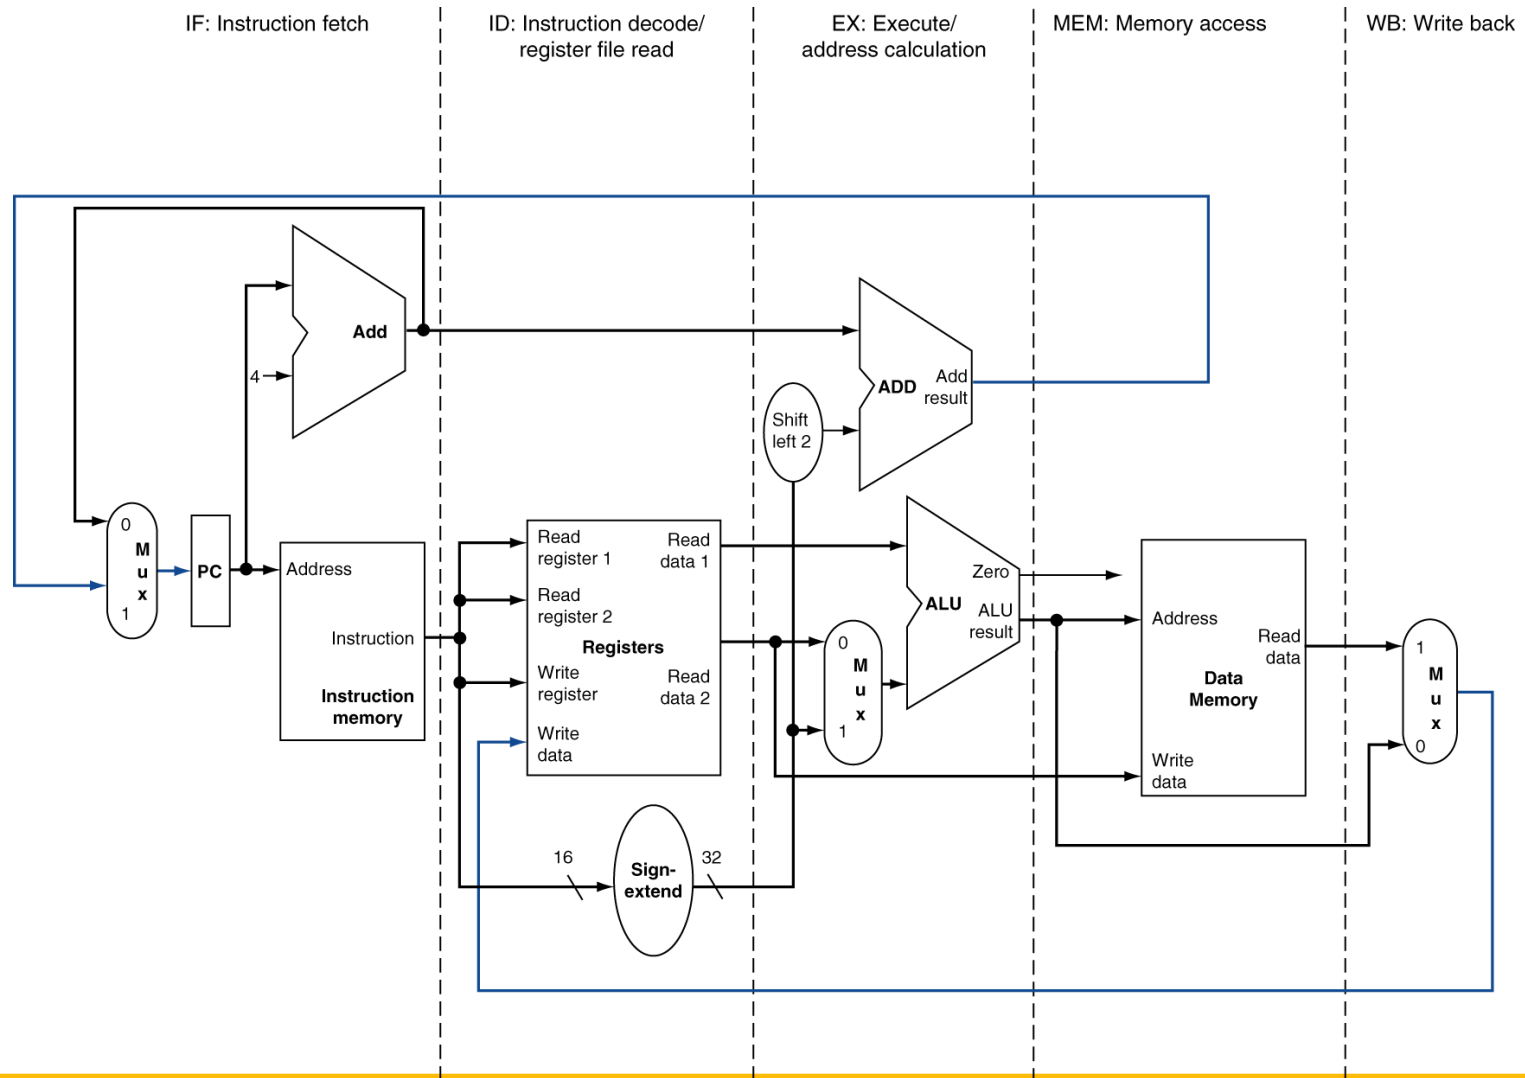
BEQ Instruction



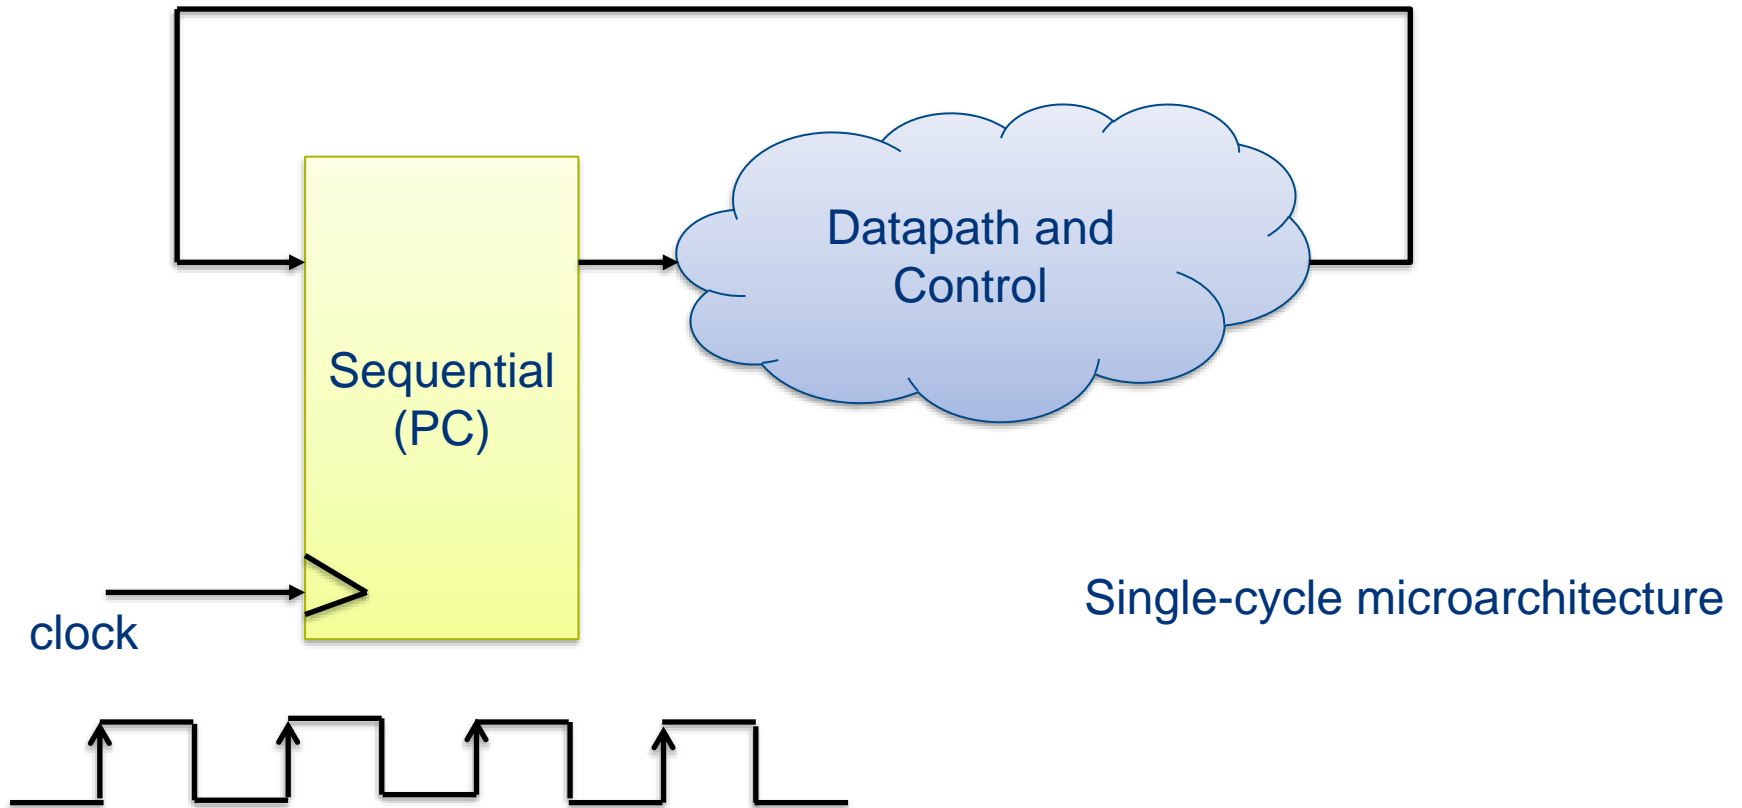
Estimating the Datapath Delay

- **Five stages**
 1. IF: Instruction fetch from memory
 2. ID: Instruction decode & register read
 3. EX: Execute operation or calculate address
 4. MEM: Access memory operand
 5. WB: Write result back to register

Five Processing Stages




Abstracting the Datapath



New instructions are fetched and executed starting at the positive edge
Each instruction must be completed within the clock period

Abstracting the Datapath



RISC-V Stages

- **R-type:**



- **I-type:**



- **Load:**



- **Store:**



- **Branch:**



- **Jump:**



Single Cycle Microarchitecture

- **Design Optimization**


- Memory units delay (read or write): 200 ps
- ALU and adders: 100 ps
- register file (read or write): 50 ps
- other combinational logic: 0 ps

steps	IF	ID	EX	MEM	WB	Delay
resources	mem	RF	ALU	mem	RF	
R-type	200	50	100		50	400
I-type	200	50	100		50	400
LW	200	50	100	200	50	600
SW	200	50	100	200		550
Branch	200	50	100			350
Jump	200					200

Single Cycle Microarchitecture

- What is the minimum clock cycle time needed?
 - 600ps

steps	IF	ID	EX	MEM	WB	Delay
resources	mem	RF	ALU	mem	RF	
R-type	200	50	100		50	400
I-type	200	50	100		50	400
LW	200	50	100	200	50	600
SW	200	50	100	200		550
Branch	200	50	100			350
Jump	200					200




Pitfalls

- **Memory access time can change**
- **Multiple memory access instructions**
- **Instructions requiring repetitions of some of the units**
- **Suggested readings:**
 - Johnson et al. “Using Multiple Memory Access Instructions for Reducing Code Size” in Compiler Construction 2004

Variable Memory Access Time

- Thus far we have assumed that data is always available in a monolithic memory
 - Processors implement “memory hierarchy”
 - Depending on where is your data the access time can change



Multiple Memory Access Instructions

- Remember addressing modes?

- Absolute LW rt, 1000
use 1000 as address
- Register Indirect: LW rt, (r_{base})
use $GPR[r_{base}]$ as address
- Displaced or base: LW rt, offset(r_{base})
use offset+ $GPR[r_{base}]$ as address
- Indexed: LW rt, (r_{base}, r_{index})
use $GPR[r_{base}] + GPR[r_{index}]$ as address
- Memory Indirect LW rt ((r_{base}))
use value at $M[GPR[r_{base}]]$ as address
- Auto increment/decrement LW Rt, (r_{base})
use $GPR[r_{base}]$ as address, but increment or decrement $GPR[r_{base}]$ each time

Single Cycle Uarch: Design Optimization

- **Thus far we have considered simple instructions**
 - In x86, many complex instructions may require multiple execution and memory units
 - True even for MIPS and ARM



- Memory units delay (read or write): 200 ps
- ALU and adders: 100 ps
- register file (read or write): 50 ps
- other combinational logic: 0 ps
- Instruction delay = $200 + 50 + 100 + 100 + 200 + 200 + 50 = 900\text{ps}$

Single Cycle Uarch: Design Optimization

- We have the scenario (common instructions)

steps	IF	ID	EX	MEM	WB	Delay
resources	mem	RF	ALU	mem	RF	
R-type	200	50	100		50	400
I-type	200	50	100		50	400
LW	200	50	100	200	50	600
SW	200	50	100	200		550
Branch	200	50	100			350
Jump	200					200

- Some special instructions requiring 900ps
- Question: should we optimize for the common case (600ps) or the worst case (900ps)?
 - Worst case. Why?
 - Otherwise, we will have design rule violation for the special cases
 - Clock is 600ps and delay is 900ps

Single Cycle Uarch: Design Optimization

- For the example above the single-cycle microarchitecture looks like



- Common case (with 5 stages)



- Special cases (with 2 EX and 2 MEM units)



Single Cycle Microarchitecture

- **Performance**

- $CPI = 1$
- Clock Cycle Time = 900 ps (worst case)
- Execution time of 100 billion instructions
 - = {# of instructions} x {Average CPI} x {clock cycle time}
 - = $\{100 \times 10^9\} \times 1 \times \{900 \times 10^{-12}\}$
 - = 90 seconds
- Average latency of each instruction = 1 cycle = 900 ps
- Average instruction throughput = $1 / (1 \times 900 \times 10^{-12}) = 1.1$ billion instructions per second
- $IPC = 1 / CPI = 1$

Performance Comparison

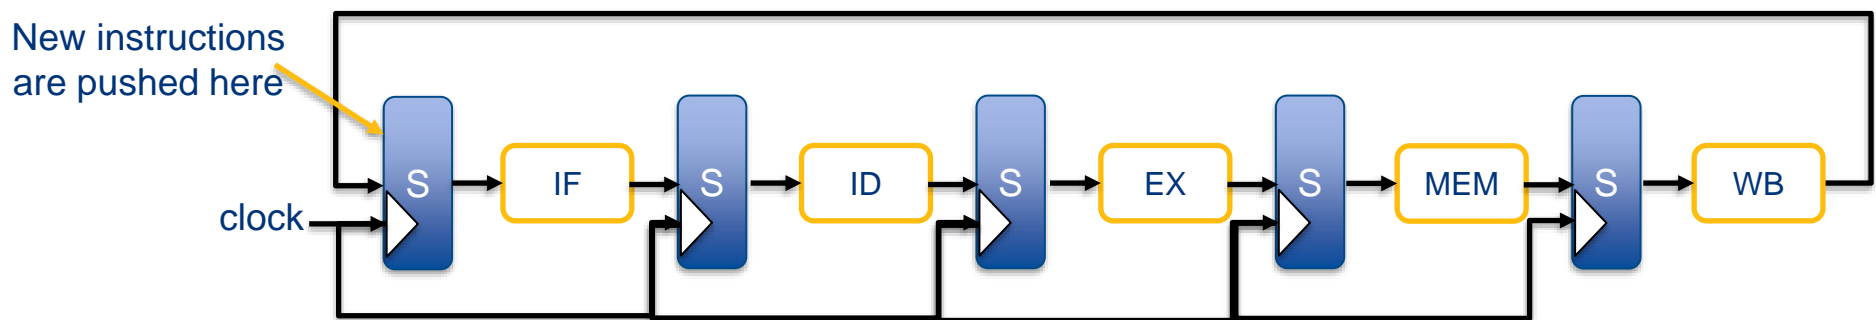
	Single cycle		
CPI	1		
Clock cycle time	900 ps		
Execution time of 100 billion instructions	90 s		
Average latency of each instruction	900 ps		
Average instruction throughput	1.1 bips		
IPC	1		

Single Cycle Microarchitecture: Concluding Remarks

- **Contrived**
 - All instructions run as slow as the slowest **instruction**
- **Inefficient**
 - All instructions run as slow as the slowest instruction
 - Must provide worst-case combinational resources in parallel as required by any instruction
 - Need to **replicate** a resource if it is needed more than once by an instruction during different parts of the instruction processing cycle
- **Not necessarily the simplest way to implement an ISA**
- **Not easy to optimize/improve performance**
 - Optimizing the common case does not work (e.g. common instructions)
 - Need to optimize the worst case all the time

Multi-cycle Microarchitecture

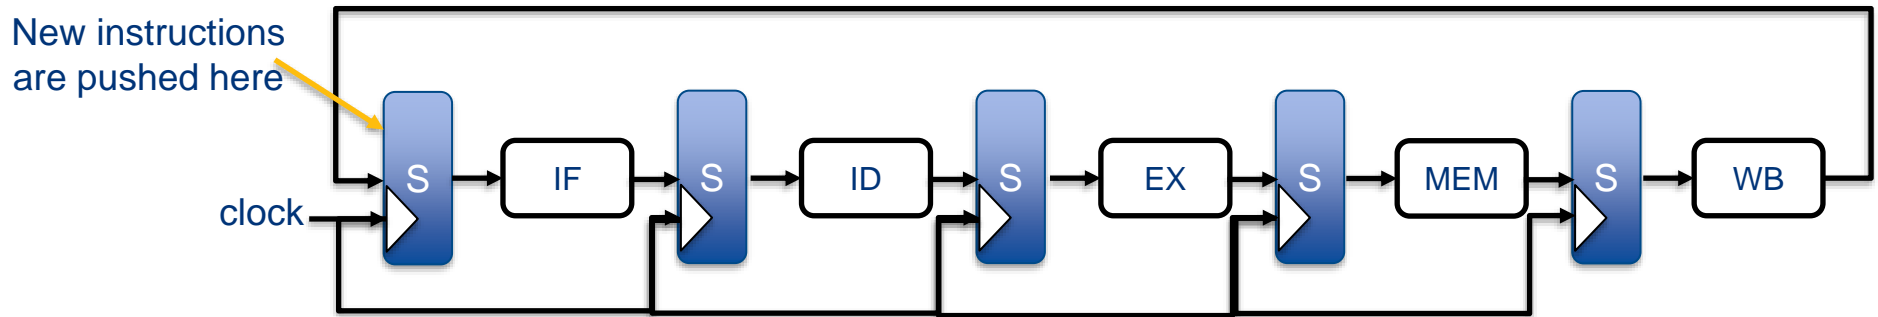
- **Idea: Start with the 5 stages and insert sequential logic after each stage**



- **How does this microarchitecture mitigate the limitations of the single-cycle architecture?**
- **How does this microarchitecture handle special long instructions?**
- **What are performance metrics using this microarchitecture?**

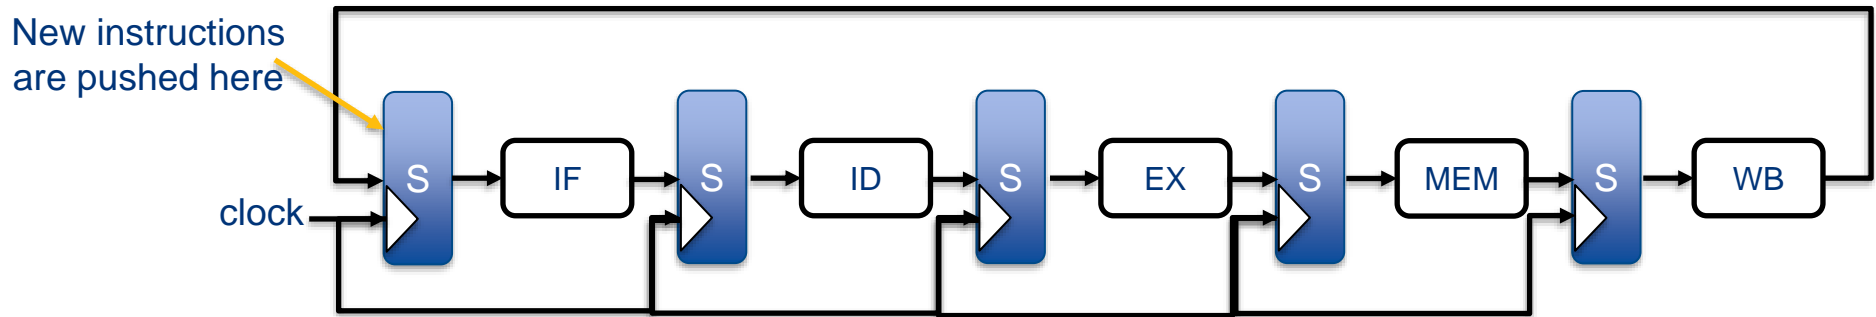
Visualization

- Inst 1
- Inst 2
- Inst 3
- ...



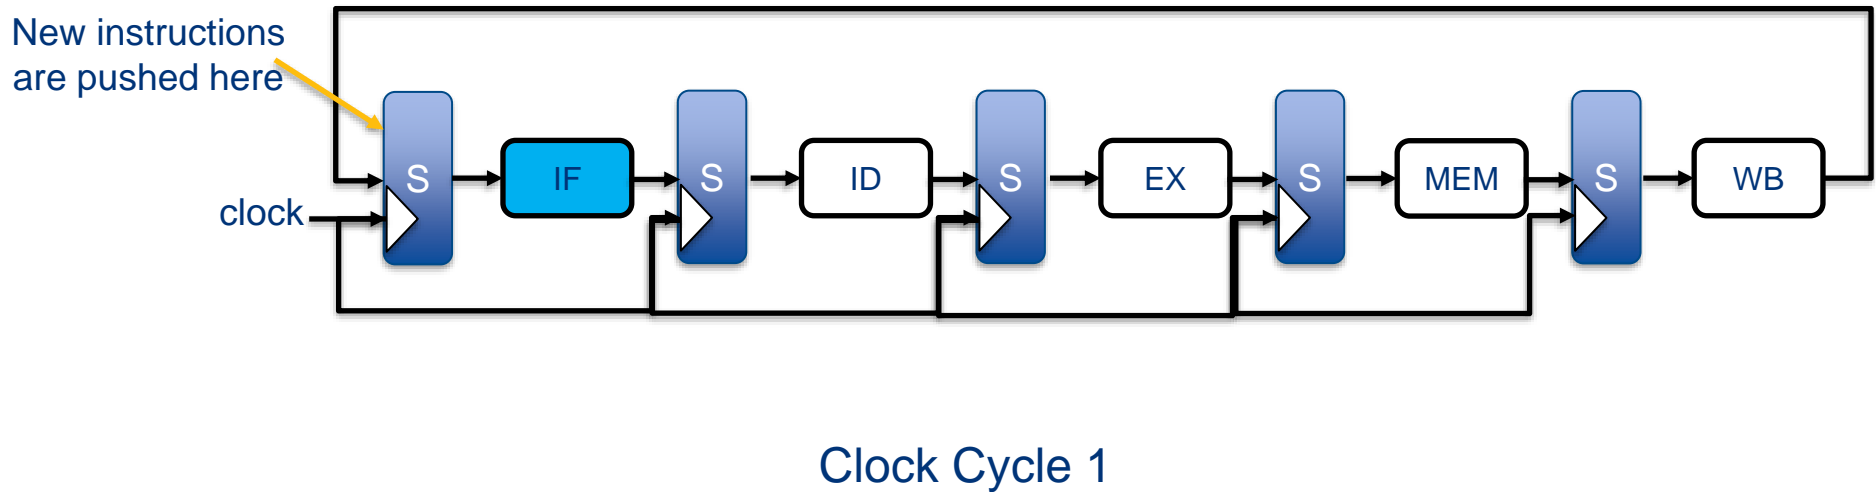
Visualization

- Inst 1 ←
- Inst 2
- Inst 3
- ...



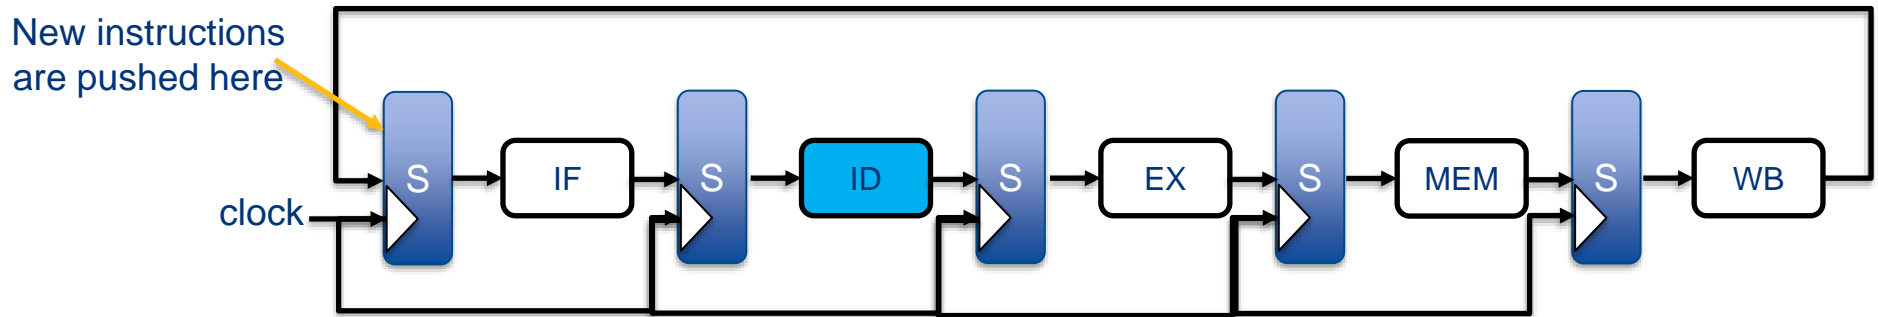
Visualization

- Inst 1 ←
- Inst 2
- Inst 3
- ...



Visualization

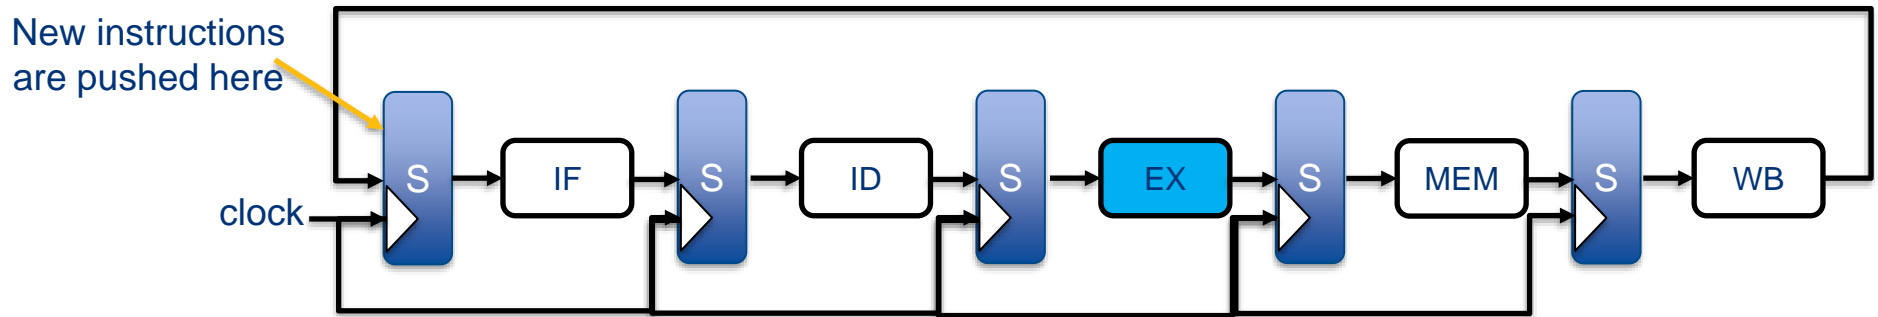
- Inst 1 ←
- Inst 2
- Inst 3
- ...



Clock Cycle 2

Visualization

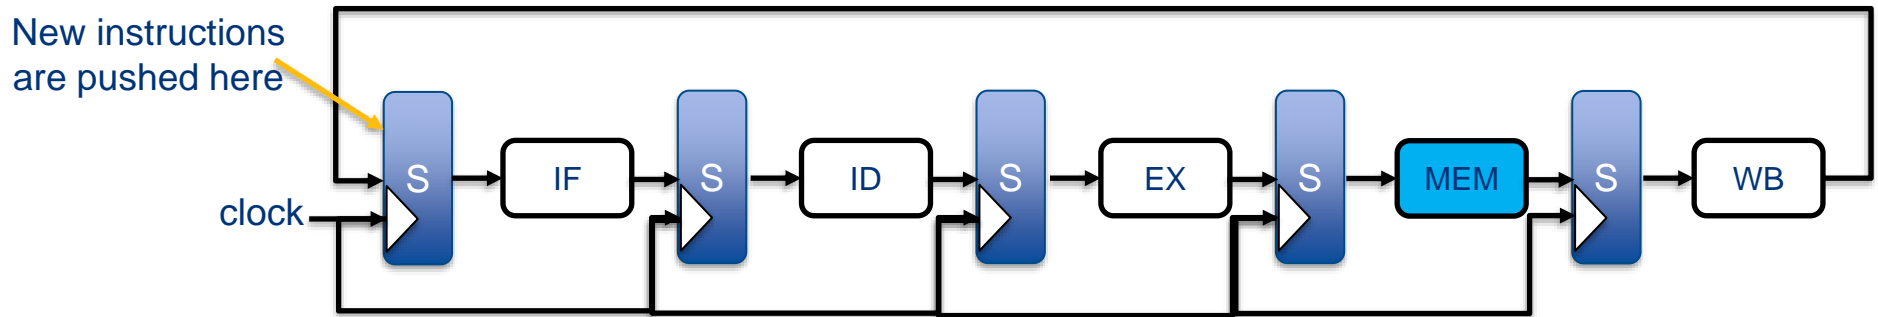
- Inst 1 ←
- Inst 2
- Inst 3
- ...



Clock Cycle 3

Visualization

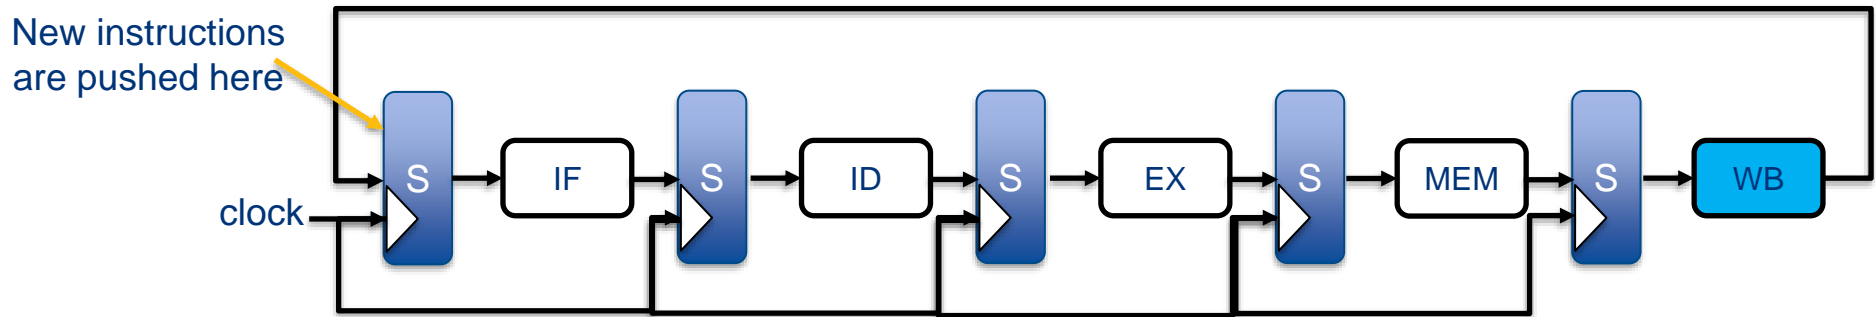
- Inst 1 ←
- Inst 2
- Inst 3
- ...



Clock Cycle 4

Visualization

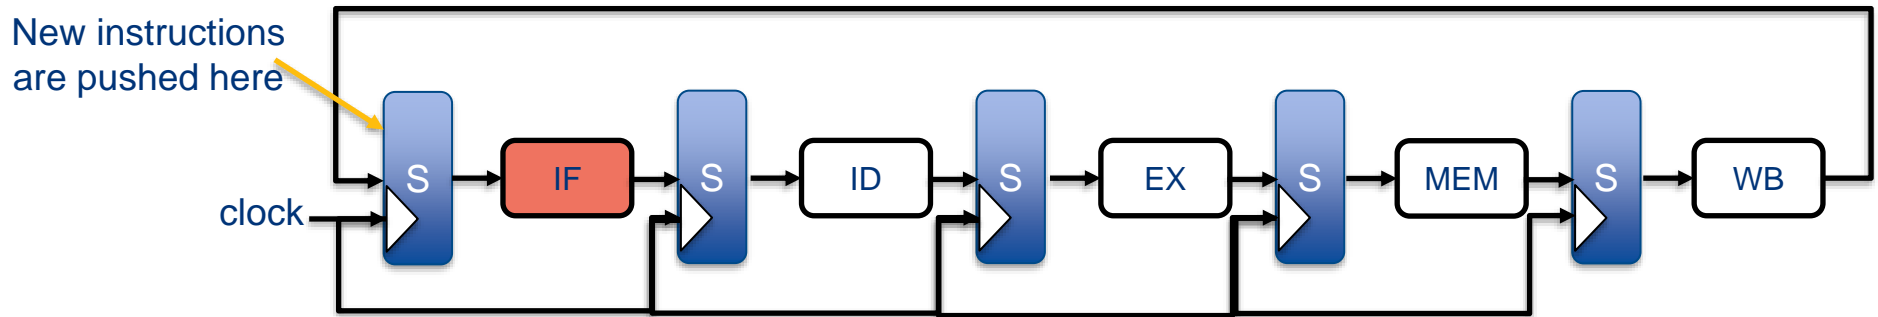
- Inst 1 ←
- Inst 2
- Inst 3
- ...



Clock Cycle 5

Visualization

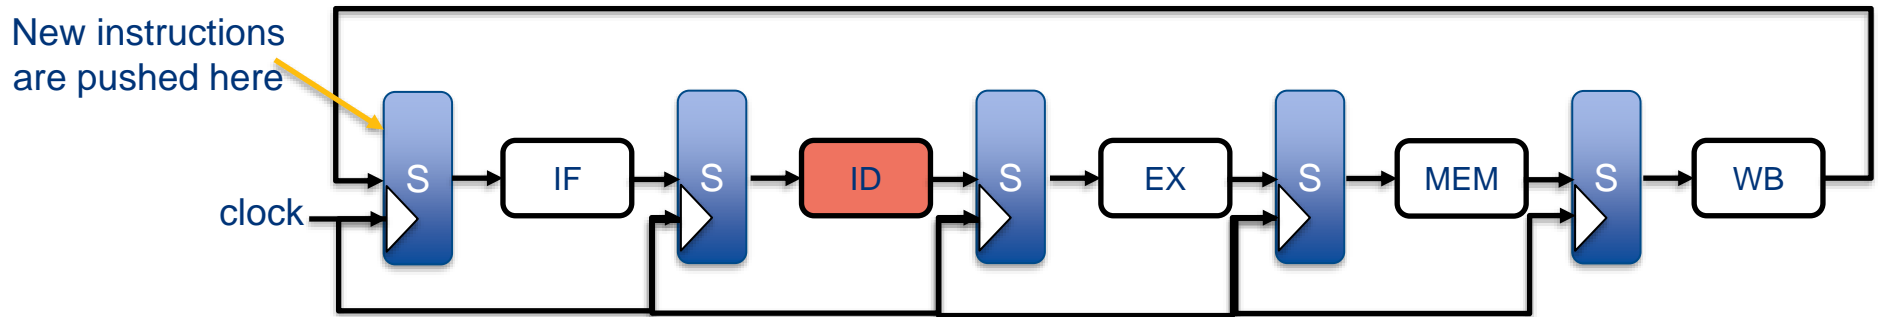
- Inst 1
- Inst 2 ←
- Inst 3
- ...



Clock Cycle 6

Visualization

- Inst 1
- Inst 2 ←
- Inst 3
- ...



Clock Cycle 7

And so on

Summary

Instr	Clocks	IF	ID	EX	MEM	WB
1	1	X				
	2		X			
	3			X		
	4				X	
	5					X
2	6	X				
	7		X			
	8			X		
	9				X	
	10					X
3	11	X				
	12		X			

Many Observations

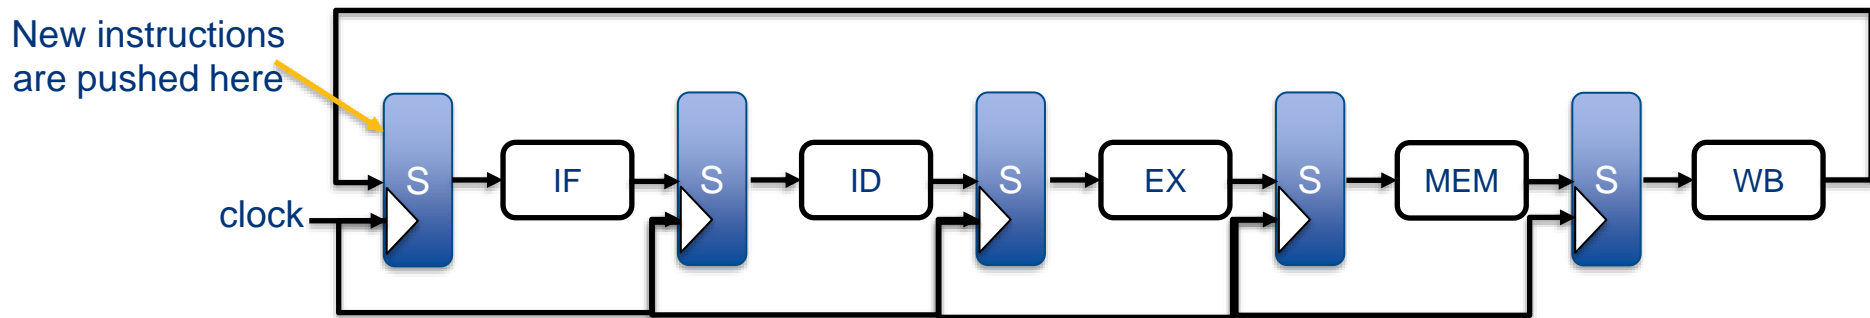
- **A new instruction starts processing only when previous instruction finishes**
- **For a given instruction, only one execution unit is occupied at any given time**
 - Under utilization of resources
 - When instruction 1 is in the decode stage, a new instruction could have used the fetch stage (pipelining, more later)
- **A single instruction now takes 5 cycles**
 - What about non standard instructions ?
(standard = one that takes all 5 stages)

Instr	Clocks	IF	ID	EX	MEM	WB
1	1	X				
	2		X			
	3			X		
	4				X	
	5					X
2	6	X				
	7		X			
	8			X		
	9				X	
	10					X
3	11	X				
	12		X			

Smaller Instructions

- **Bypass stages**

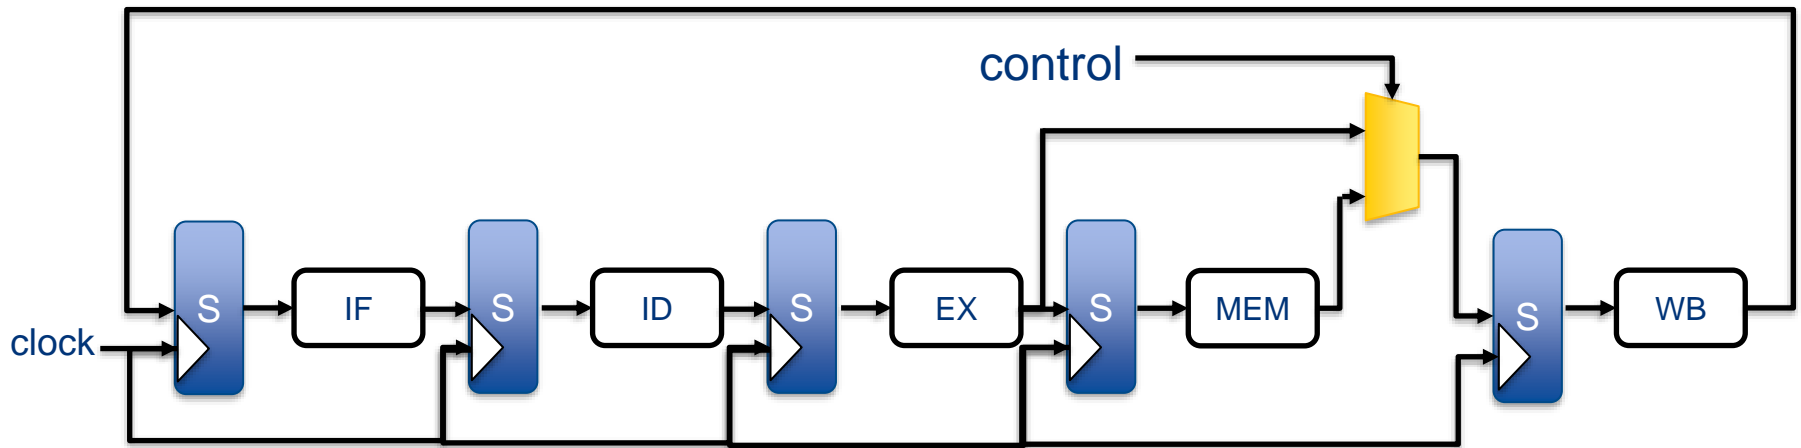
- Example of R-type Instruction: IF, ID, EX, WB (no MEM stage needed)



Smaller Instructions

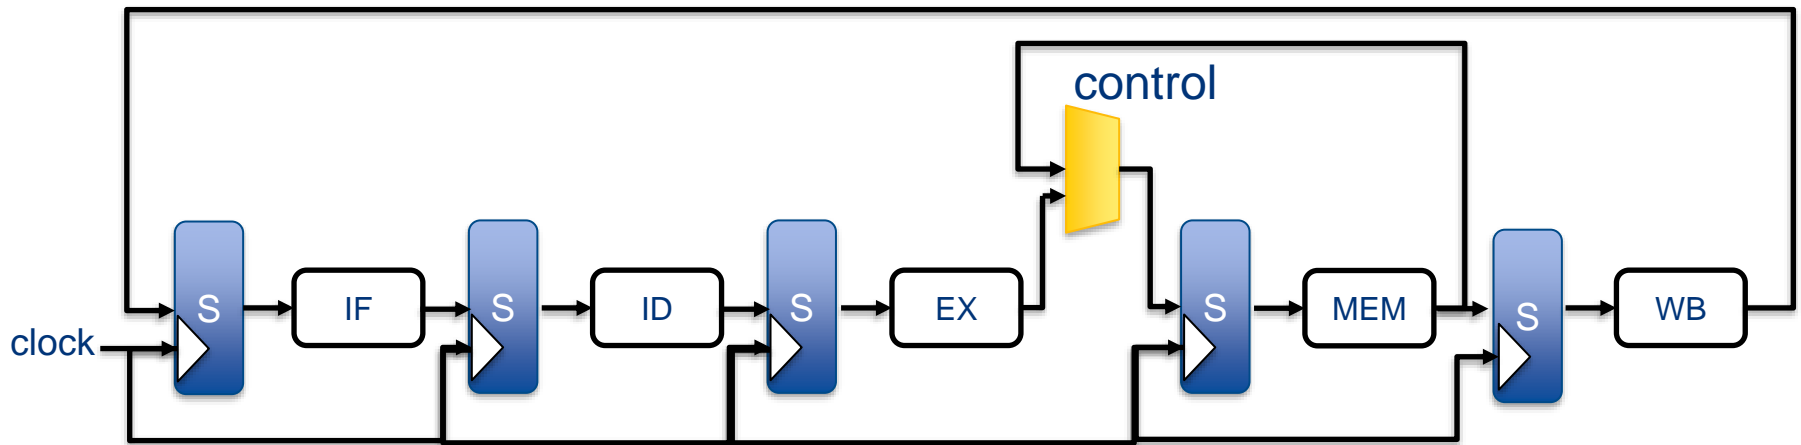
- **Bypass stages**

- Example of R-type Instruction: IF, ID, EX, WB (no MEM stage needed)



Longer Instructions

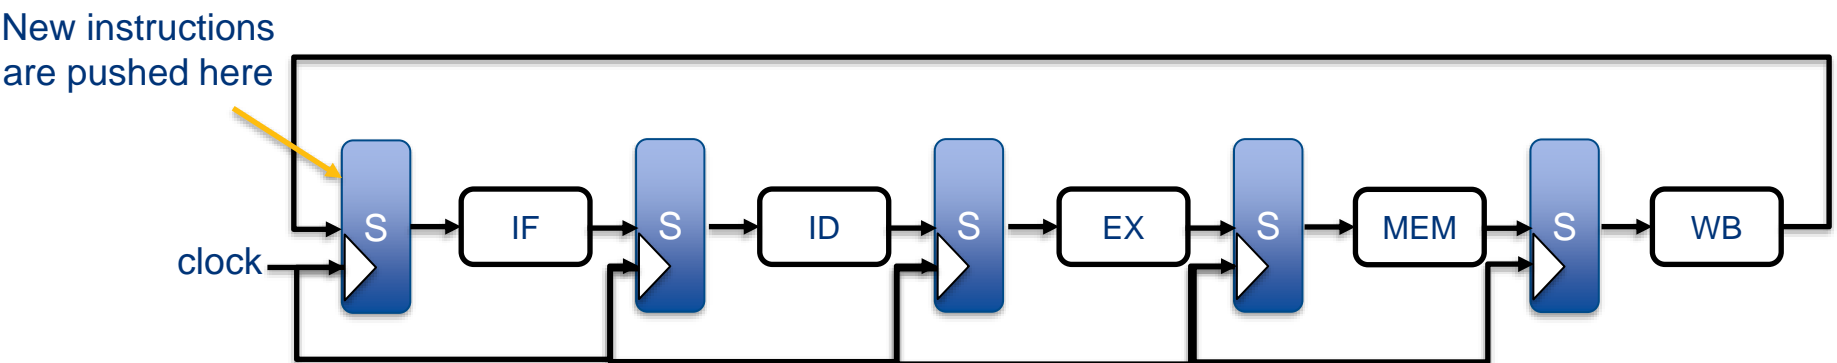

- **Example: Instruction requiring 2 MEM stages**



Multi-cycle Microarchitecture

- **Singe vs Mult-Cycle**

Critical path = entire stages



Critical path = each stage

Multi-cycle Microarchitecture

- **Critical path design**

- Can keep reducing the critical path independently of the worst-case processing time of any instruction

- **common case design**

- Can optimize the number of states it takes to execute “important” instructions that make up much of the execution time

- **Balanced design**

- No need to provide more capability or resources than really needed
 - An instruction that needs resource X multiple times does not require multiple X's to be implemented
 - Leads to more efficient hardware: Can reuse hardware components needed multiple times for an instruction

Multi-cycle Microarchitecture: performance

- **Execution time of an instruction**
 - $\{\text{CPI}\} \times \{\text{clock cycle time}\}$
- **Execution time of a program**
 - Sum over all instructions $[\{\text{CPI}\} \times \{\text{clock cycle time}\}]$
 - $\{\# \text{ of instructions}\} \times \{\text{Average CPI}\} \times \{\text{clock cycle time}\}$
- **Single cycle microarchitecture performance**
 - $\text{CPI} = 1$
 - Clock cycle time = long
- **Multi-cycle microarchitecture performance**
 - $\text{CPI} = \text{different for each instruction}$
 - Average CPI \rightarrow hopefully small
 - Clock cycle time = short

Multi-cycle Microarchitecture: performance

- **Instruction takes different number of cycles**
 - 3 cycles
 - branch
 - 4 cycles
 - R-type, store, addi
 - 5 cycles
 - load
- **CPI = Average/weighted average**
 - SPECINT2000 benchmarks
 - 25% load
 - 10% store
 - 13% branches
 - 52% R-type
 - Average CPI = $(25*5 + 10*4 + 13*3 + 52*4) / (25 + 10 + 13 + 52)$
 $= 4.12$

Multi-cycle Microarchitecture: Performance

steps	IF	ID	EX	MEM	WB	Delay
resources	mem	RF	ALU	mem	RF	
R-type	200	50	100		50	400
I-type	200	50	100		50	400
LW	200	50	100	200	50	600
SW	200	50	100	200		550
Branch	200	50	100			350
Jump	200					200

What should be the clock cycle?: 3 numbers: 200, 100, 50
- 200ps

Multi-cycle Microarchitecture: Performance

- **CPI = 4.12**
- **Clock Cycle Time = 200ps**
- **Execution time of 100 billion instructions**
 - $= \{\text{\# of instructions}\} \times \{\text{Average CPI}\} \times \{\text{clock cycle time}\}$
 - $= \{100 \times 10^9\} \times 4.12 \times \{200 \times 10^{-12}\}$
 - = 82.4 seconds
- **Average latency of each instruction = 4.12 cycles = 4.12 x 200 = 824ps**
- **Average instruction throughput = $1 / (4.12 \times 200 \times 10^{-12})$ = 1.21 billion instructions per second**
- **IPC = $1 / \text{CPI} = 1 / 4.12 = 0.24$**

Performance Comparison

	Single cycle	Multi-cycle	
CPI	1		
Clock cycle time	900 ps		
Execution time of 100 billion instructions	90 s		
Average latency of each instruction	900 ps		
Average instruction throughput	1.1 bips		

Performance Comparison

	Single cycle	Multi-cycle	
CPI	1	4.12	
Clock cycle time	900ps	200ps	
Execution time of 100 billion instructions	90s	82.4s	
Average latency of each instruction	900ps	824ps	
Average instruction throughput	1.1 bips	1.21 bips	

How can we improve these metrics?

Representations

clock State of the microarchitecture

1

F	D	E	M	W
---	---	---	---	---

Representations

clock State of the microarchitecture

1

F	D	E	M	W
---	---	---	---	---

2

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W
3					

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W
3	F	D	E	M	W

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W
3	F	D	E	M	W
4					

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W
3	F	D	E	M	W
4	F	D	E	M	W

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W
3	F	D	E	M	W
4	F	D	E	M	W
5					

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W
3	F	D	E	M	W
4	F	D	E	M	W
5	F	D	E	M	W

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W
3	F	D	E	M	W
4	F	D	E	M	W
5	F	D	E	M	W
6					

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W
3	F	D	E	M	W
4	F	D	E	M	W
5	F	D	E	M	W
6	F	D	E	M	W

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W
3	F	D	E	M	W
4	F	D	E	M	W
5	F	D	E	M	W
6	F	D	E	M	W
7					

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W
3	F	D	E	M	W
4	F	D	E	M	W
5	F	D	E	M	W
6	F	D	E	M	W
7	F	D	E	M	W

Representations

clock State of the microarchitecture

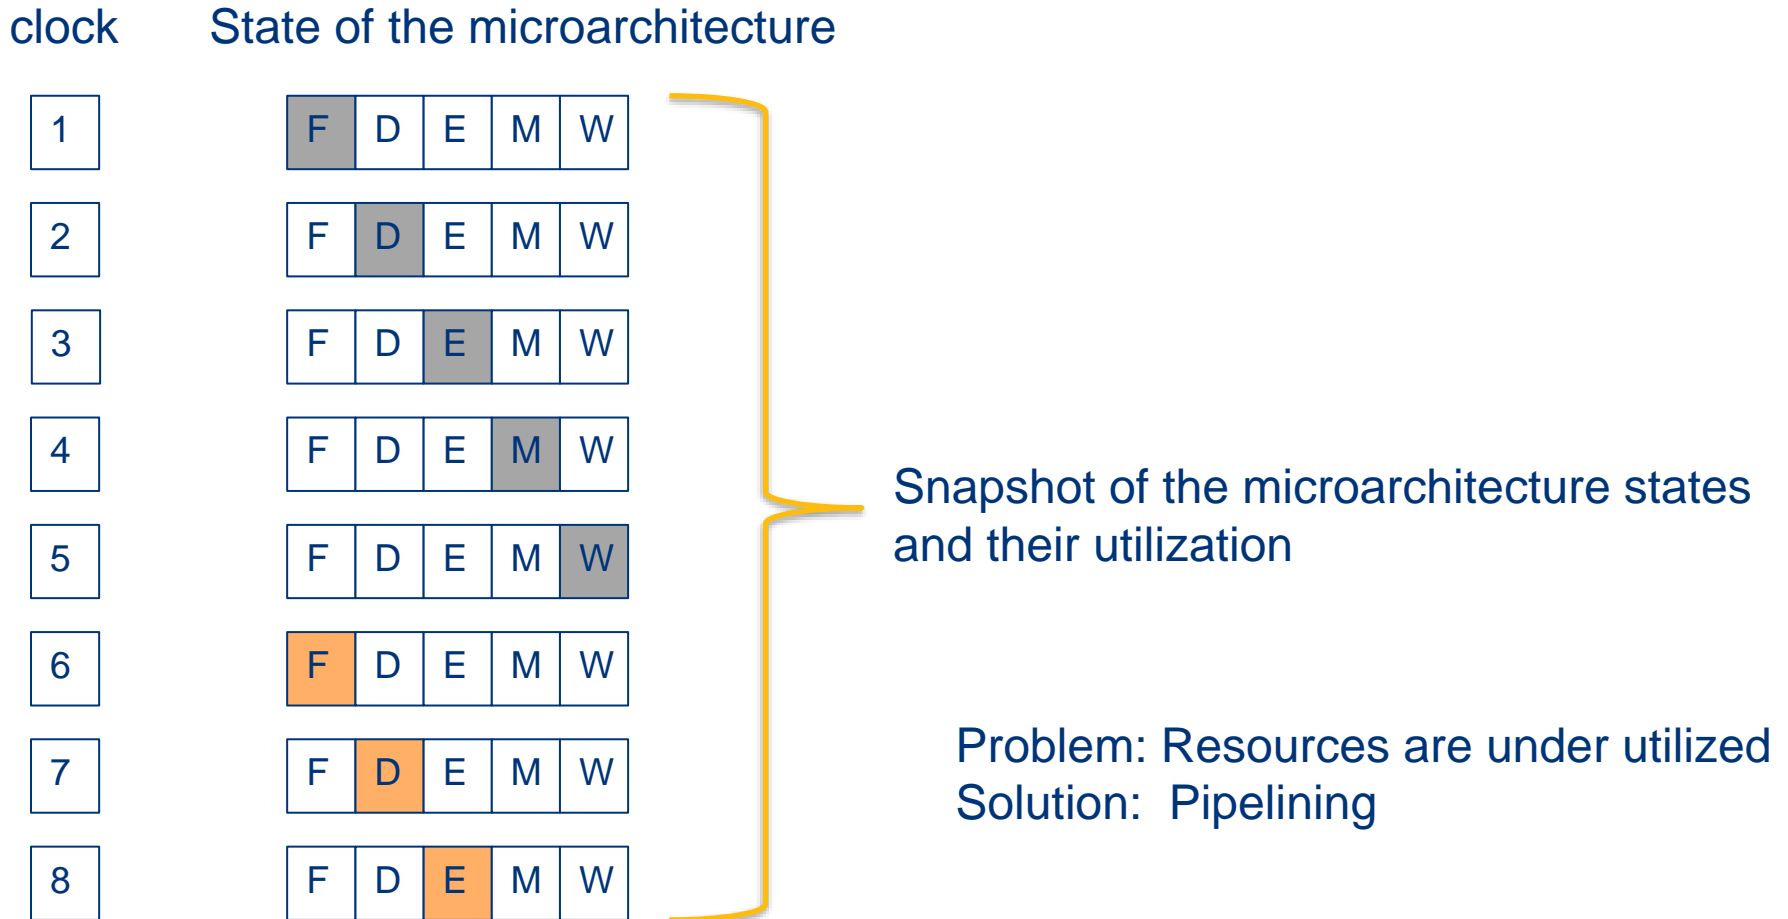
1	F	D	E	M	W
2	F	D	E	M	W
3	F	D	E	M	W
4	F	D	E	M	W
5	F	D	E	M	W
6	F	D	E	M	W
7	F	D	E	M	W
8					

Representations

clock State of the microarchitecture

1	F	D	E	M	W
2	F	D	E	M	W
3	F	D	E	M	W
4	F	D	E	M	W
5	F	D	E	M	W
6	F	D	E	M	W
7	F	D	E	M	W
8	F	D	E	M	W

Representations



Improve Performance of Multi-Cycle Microarchitecture


- **Can We Use the Idle Hardware to Improve Concurrency?**
 - Goal: More concurrency → Higher instruction throughput (i.e., more “work” completed in one cycle)
 - Idea: When an instruction is using some resources in its processing phase, **process other instructions on idle resources** not needed by that instruction
 - E.g., when an instruction is being decoded, fetch the next instruction
 - E.g., when an instruction is being executed, decode another instruction
 - E.g., when an instruction is accessing data memory (ld/st), execute the next instruction
 - E.g., when an instruction is writing its result into the register file, access data memory for the next instruction
- **Pipelining ☺**

Pipelining for instruction processing:

Basic Idea



- **How to achieve concurrency in instruction processing?**
 - **Divide** the instruction processing cycle into distinct “**stages**” of processing
 - Ensure there are **enough hardware resources** to process one instruction in each **stage**
 - Process a **different instruction** in each **stage**
 - Instructions consecutive in program order are processed in consecutive stages
 - Process **each stage** in a single **clock cycle**
 - Pitfalls (later)
- **Benefit: Increases instruction processing throughput (1/CPI)**
- **Drawback: ?**
- **Quick question: Is pipelining a single-cycle or a multi-cycle microarchitecture?**

The Laundry Analogy





- “place one dirty load of clothes in the washer”
 - “when the washer is finished, place the wet load in the dryer”
 - “when the dryer is finished, take out the dry load and fold”
 - “when folding is finished, put the clothes in closet”
- steps to do a load are sequentially dependent
 - no dependence between different loads
 - different steps do not share resources
 - one load every 2 hours

The Laundry Analogy



- 4 loads of laundry in parallel
- no additional resources
- latency per load is the same
- each load finishing every half hour
- throughput increased by 4

Pipelining Multiple Loads of Laundry: In Practice





- each load finishing every hour
- throughput reduced by $\frac{1}{2}$
- the slowest step decides throughput

An Ideal Pipeline

- **Increase throughput with little increase in cost**
 - hardware cost, in case of instruction processing
- **Repetition of identical operations**
 - The same operation is repeated on a large number of different inputs (e.g., all laundry loads go through the same steps)
- **Repetition of independent operations**
 - No dependencies between repeated operations
- **Uniformly partitionable suboperations**
 - Processing can be evenly divided into uniform-latency suboperations (that do not share resources)
 - In the previous example, everything takes 30 mins


Pipelining Multiple Loads of Laundry: In Practice




throughput restored (2 loads per hour) using 2 dryers

Pipelining for Instruction Processing


- **Multi-cycle: 5 cycles per instruction**



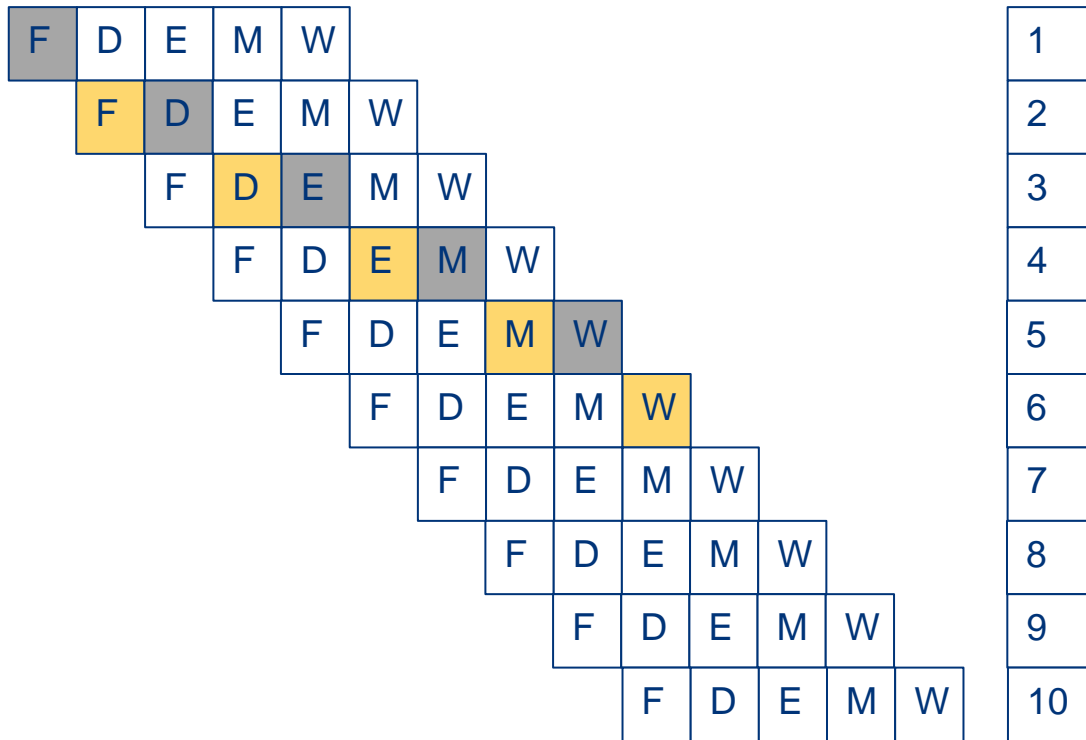
- **Pipelined:**




Example: Execution of Independent Instructions




Example: Execution of Independent Instructions




Example: Execution of Independent Instructions




Example: Execution of Independent Instructions



Example: Execution of Independent Instructions




Example: Execution of Independent Instructions



Pipelining: CPI

Example: Execution of Independent Instructions



Key Facts about Pipelining

- **There are no instructions completed for the first 4 cycles**
- **Once all the stages are working (i.e., the pipeline is full), instructions start getting completed**
 - We call the state where the pipeline becomes full as the steady state
 - IPC/CPI are computed in steady state
- **What are the steady state metrics?**
 - CPI?
 - 1
 - IPC?
 - 1
 - What about instruction latency?
 - 5 cycles
 - Does the instruction latency greater than, equal to, or lower than that of multi-cycle microarchitecture?

Revisit

- **Multi-cycled vs pipelined microarchitecture**
 - Both has multiple “**stages**” between the initial and final “**states**” which are **clocked**
 - Multi-cycled microarchitecture
 - $CPI > 1$, $IPC < 1$
 - Pipelined microarchitecture
 - $CPI = 1$ and $IPC = 1$
- **Lets keep them as separate categories**

Pipelined Microarchitecture: Performance

- **CPI = 1**
- **Clock Cycle Time = 200ps**
- **Execution time of 100 billion instructions**
 - $= \{\text{\# of instructions}\} \times \{\text{Average CPI}\} \times \{\text{clock cycle time}\}$
 - $= \{100 \times 10^9\} \times 1 \times \{200 \times 10^{-12}\}$
 - = 20 seconds
- **Average latency of each instruction = 5 cycles = $5 \times 200 = 1000\text{ps}$**
- **Average instruction throughput = $1 / (1 \times 200 \times 10^{-12}) = 5$ billion instructions per second**
- **IPC = $1 / \text{CPI} = 1$**


Recap of Multi-Cycle Microarchitecture

	Single cycle	Multi-cycle	
CPI	1	4.12	
Clock cycle time	900ps	200ps	
Execution time of 100 billion instructions	90s	82.4s	
Average latency of each instruction	900ps	824ps	
Average instruction throughput	1.1 bips	1.21 bips	
IPC	$= 1/1 = 1$	$= 1/4.12 = 0.24$	

Recap of Multi-Cycle Microarchitecture


	Single cycle	Multi-cycle	Pipeline
CPI	1	4.12	1
Clock cycle time	900ps	200ps	200ps
Execution time of 100 billion instructions	90s	82.4s	20s
Average latency of each instruction	900ps	824ps	1000ps
Average instruction throughput	1.1 bips	1.21 bips	5 bips
IPC	$1/1 = 1$	$1/4.12 = 0.24$	$1/1 = 1$

Example: Execution of Independent Instructions




Only a part of the instruction gets processed at any given clock cycle
Different stages process different instructions (part of it) at any given clock cycle


Example: Execution of Independent Instructions




Example: Execution of Instructions



Example: Execution of Instructions




Example: Execution of Instructions



What does the programmer see?


- **Expected Instruction Order**

- **Obtained Instruction Order**




Do you see a difference in instruction order?
Not acceptable!


Example: Execution of Instructions




Example: Execution of Instructions



Example: Execution of Instructions




Example: Execution of Instructions



Incorrect Execution

- **Discarding all instructions that are currently in the pipeline**
 - Also called **pipeline flushing**
 - All stages of the pipeline are processing some instruction
 - Important: Only those instructions are flushed which are in and prior to the stage where error happened!!
 - How many total instructions discarded?
 - 3
 - On what this number “3” depends on?
 - The stage where error happened
 - If the error happened in the decode stage, how many instructions will be flushed?
 - 2

Example: Execution of Instructions



Incorrect Execution

- **Discarding all instructions that are currently in the pipeline**
 - Look at the time to process instruction. How many cycles are there where no instruction is completed?
 - 3
 - On what does it depends?
 - The stage where error happened
 - All cycles where there is no valid instruction output, we call **stall cycles** or **bubbles**
 - How many stall cycles?
 - 3

Impact of Bubbles in pipeline



- **100 instructions**
 - No bubbles
 - 100 instructions in 100 cycles (steady state)
 - $\text{CPI} = 100/100 = 1$
 - 1 instruction in error (3 bubbles)
 - 100 instructions in 103 cycles (steady state)
 - $\text{CPI} = 103/100 = 1.03$
 - 10 instructions in error (3 bubbles/incorrect instruction)
 - 100 instructions in 130 cycles (steady state)
 - $\text{CPI} = 130/100 = 1.3$
- **Observe how CPI is climbing up**
 - High CPI is bad

Pipeline bubbles concluding remarks

- Pipeline bubbles are bad and must be avoided
- Bubbles once introduced, stays in the pipeline



.

The Laundry Analogy





- 4 loads of laundry in parallel
- no additional resources
- latency per load is the same
- each load finishing every half hour
- throughput increased by 4

Pipelining Multiple Loads of Laundry: In Practice



- each load finishing every half hour




- each load finishing every hour

- throughput reduced by $\frac{1}{2}$

- the slowest step decides throughput

Theoretical Analysis

- 5-stage pipeline



Throughput = 1 / delay of stage


Theoretical Analysis

- 2-stage pipeline

Logic

Theoretical Analysis

- 2-stage pipeline



Throughput = $1 / \text{delay of stage}$


Theoretical Analysis

- 3-stage pipeline

Logic

Theoretical Analysis


- 3-stage pipeline



Throughput = $1 / \text{delay of stage}$

Theoretical Analysis

- 3-stage pipeline



Throughput = 1 / delay of stage

Theoretical Analysis

- 4-stage pipeline

Logic

Theoretical Analysis

- 4-stage pipeline


Logic

M	N	O	P
---	---	---	---

Throughput = $1 / \text{delay of stage}$

Theoretical Analysis


- 4-stage pipeline




Throughput = 1 / delay of stage

Delay of a stage


- **2-stage pipeline**



- **3-stage pipeline**



- **4-stage pipeline**



Delay of a stage

- **2-stage pipeline**



Delay of a stage = $T/2$

- **3-stage pipeline**



Delay of a stage = $T/3$

- **4-stage pipeline**



Delay of a stage = $T/4$

Delay of a stage

- 2-stage pipeline

 Latch with delay S



Delay of a stage = $T/2 + S$



- 3-stage pipeline



Delay of a stage = $T/3 + S$



- 4-stage pipeline



Delay of a stage = $T/4 + S$



Delay of a stage

- 2-stage pipeline

 Latch with delay S



$$\text{Delay of a stage} = T/2 + S$$



- 3-stage pipeline



$$\text{Delay of a stage} = T/3 + S$$



$$\text{Throughput} = 1 / \text{delay}$$

- 4-stage pipeline



$$\text{Delay of a stage} = T/4 + S$$



Delay of a stage

- 2-stage pipeline

 Latch with delay S



Delay of a stage = $T/2 + S$ Throughput = $1 / (T/2 + S)$

- 3-stage pipeline



Delay of a stage = $T/3 + S$ Throughput = $1 / (T/3 + S)$

- 4-stage pipeline



Delay of a stage = $T/4 + S$ Throughput = $1 / (T/4 + S)$

Delay of a stage

- 2-stage pipeline

 Latch with delay S



Delay of a stage = $T/2 + S$ **Throughput = $1 / (T/2 + S)$**

- 3-stage pipeline



Delay of a stage = $T/3 + S$ **Throughput = $1 / (T/3 + S)$**

- 4-stage pipeline



Delay of a stage = $T/4 + S$ **Throughput = $1 / (T/4 + S)$**

Throughput = $1 / (T/k + S)$, k = number of stages

Area Analysis

- 2-stage pipeline

 Latch with area L



$$A = G/2, B = G/2$$

$$\text{Total area} = G/2 + L + G/2 + L = \mathbf{G + 2L}$$

- 3-stage pipeline



$$X = G/3, Y = G/3, Z = G/3$$

$$\text{Total area} = G/3 + L + G/3 + L + G/3 + L = \mathbf{G + 3L}$$

- 4-stage pipeline



$$M = G/4, N = G/4, O = G/4, P = G/4$$

$$\text{Total area} = \mathbf{G + 4L}$$

Area Analysis

- 2-stage pipeline

 Latch with area L



$$A = G/2, B = G/2$$

$$\text{Total area} = G/2 + L + G/2 + L = G + 2L$$

$$\text{Area} = G + k \cdot L, k = \text{number of stages}$$

- 3-stage pipeline



$$X = G/3, Y = G/3, Z = G/3$$

$$\text{Total area} = G/3 + L + G/3 + L + G/3 + L = G + 3L$$

- 4-stage pipeline



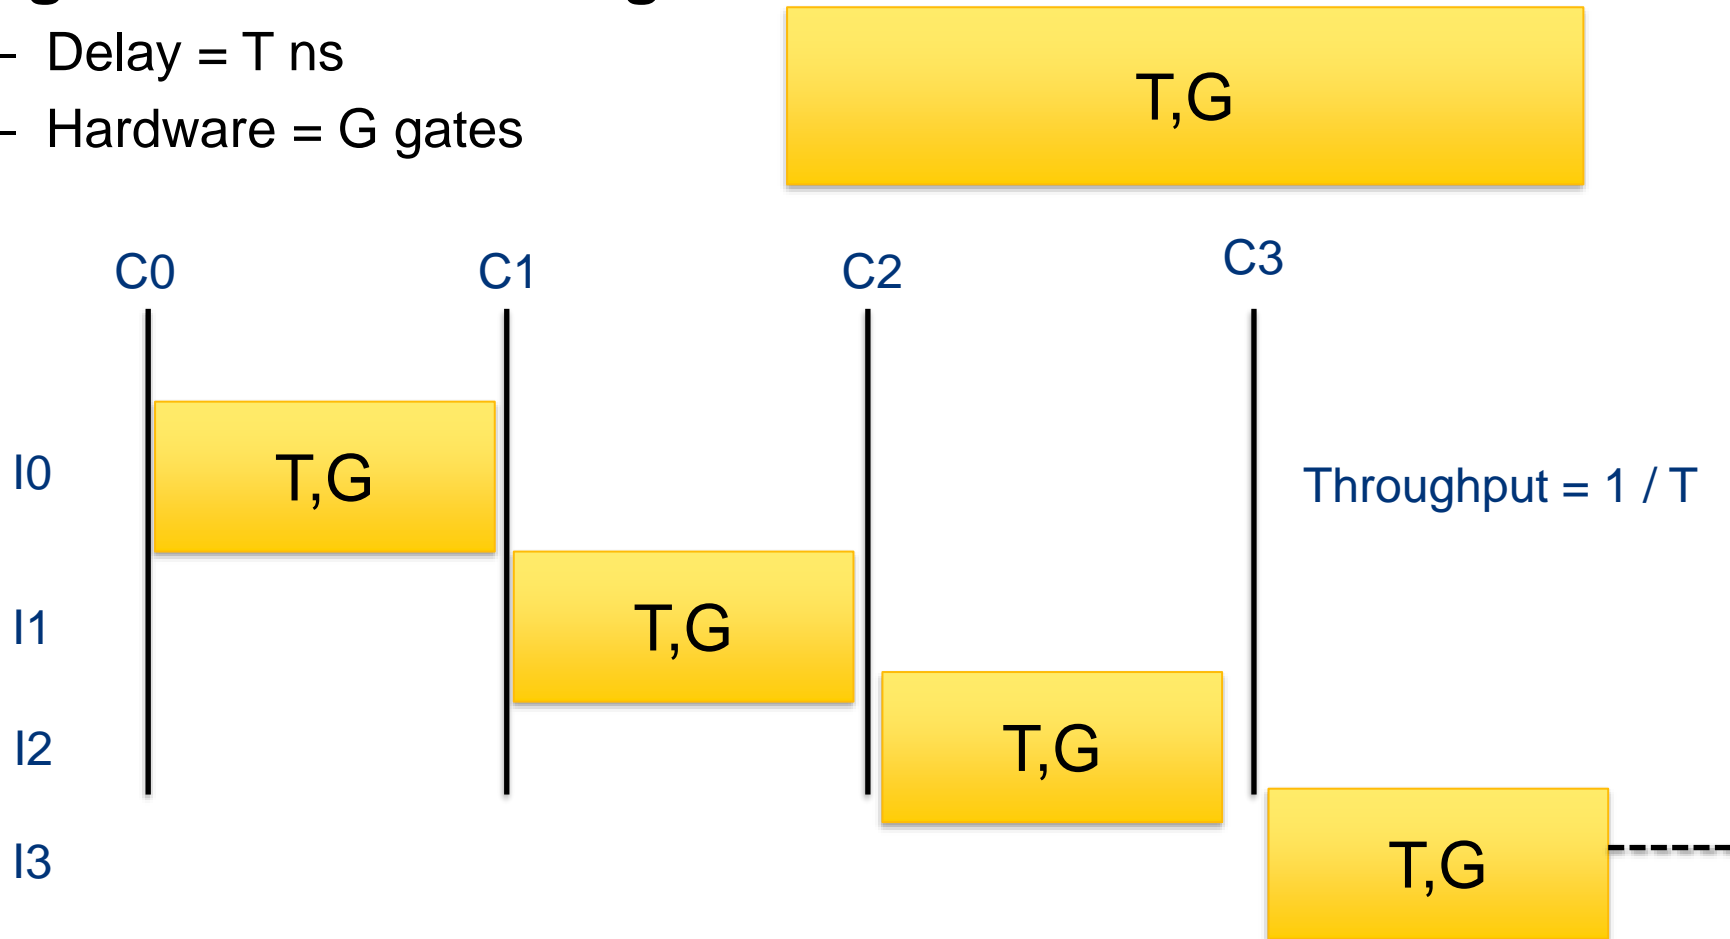
$$M = G/4, N = G/4, O = G/4, P = G/4$$

$$\text{Total area} = G + 4L$$

Theoretical Analysis

- **Big Combinational Logic**

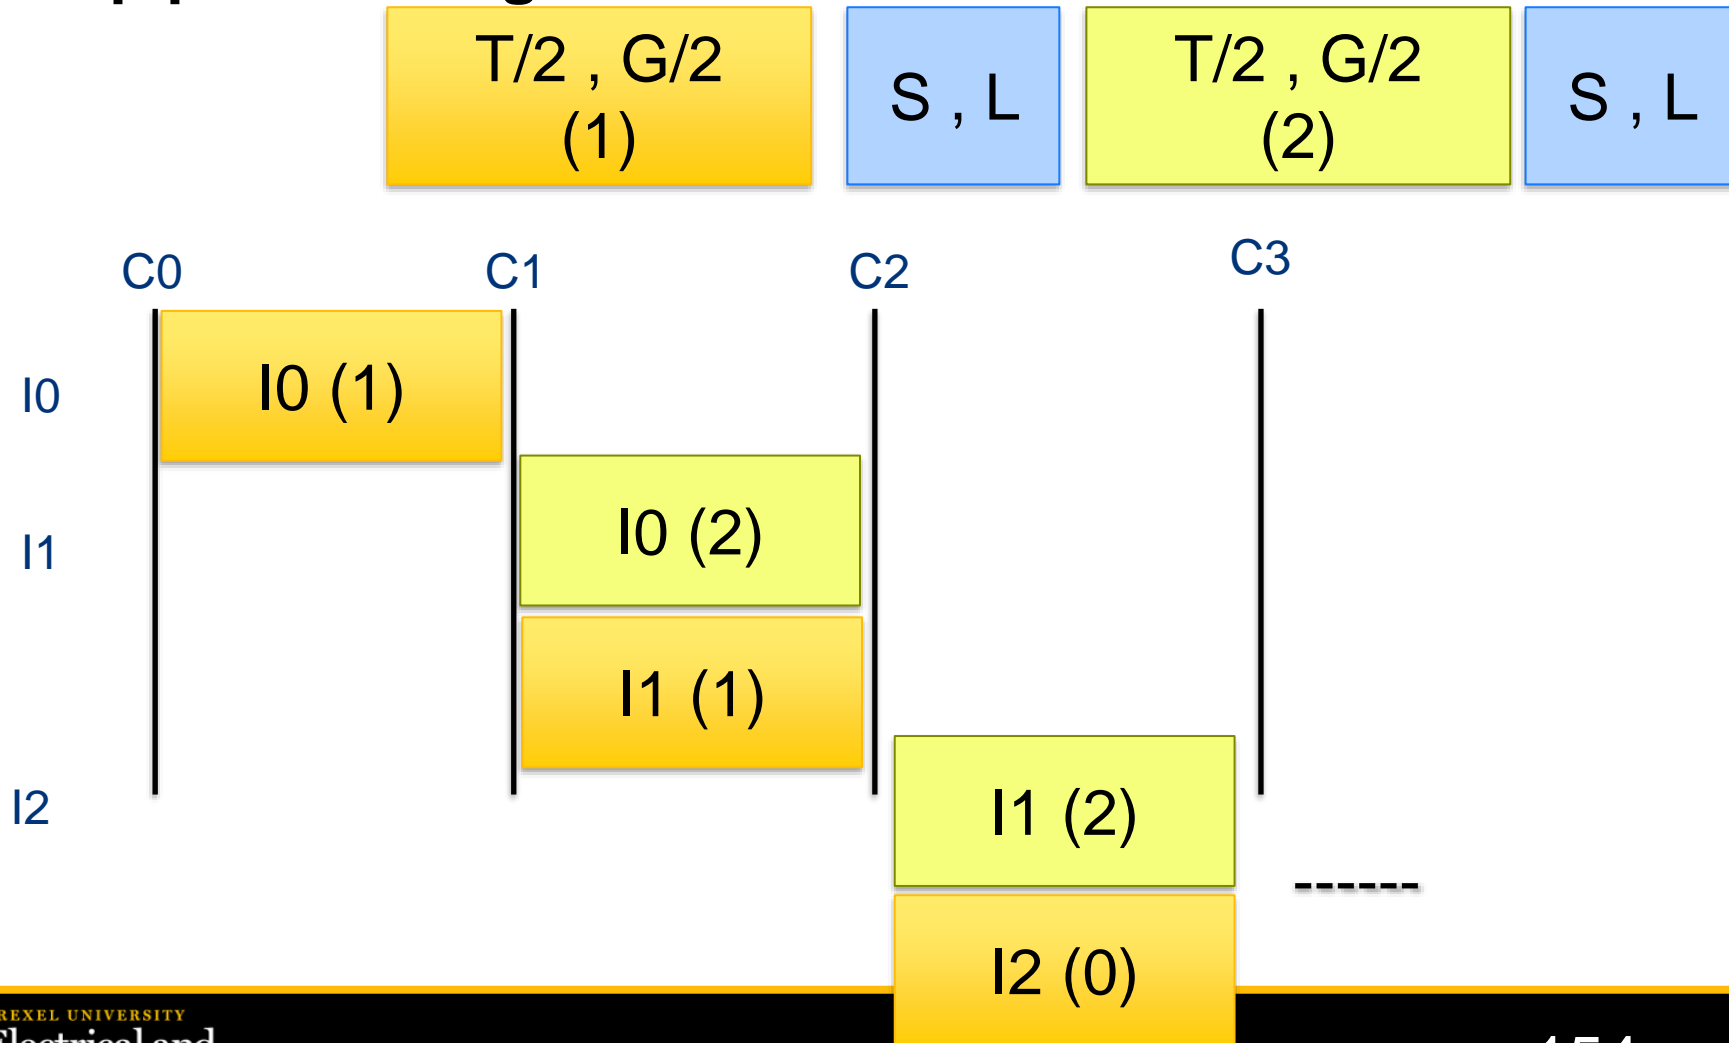
- Delay = T ns
- Hardware = G gates



Theoretical Analysis

$$\text{Throughput} = 1 / (T/2 + S)$$

- Two pipelined stages




Theoretical Analysis

- **Throughput = $1 / (T/2 + S)$**
- **Area = $G/2 + L + G/2 + L = G + 2*L$**

Theoretical Analysis

- For 3 stages



- Throughput = $1 / \text{delay of a stage} = 1 / (T/3 + S)$
- Total Area = $G/3 + L + G/3 + L + G/3 + L = G + 3*L$

Theoretical Analysis

- **k-stages pipeline**
- **Throughput = $1 / \text{delay of a stage} = 1 / (T/k + S)$**
- **Area = $G + k \cdot L$**

More Realistic Pipeline: Throughput and Cost

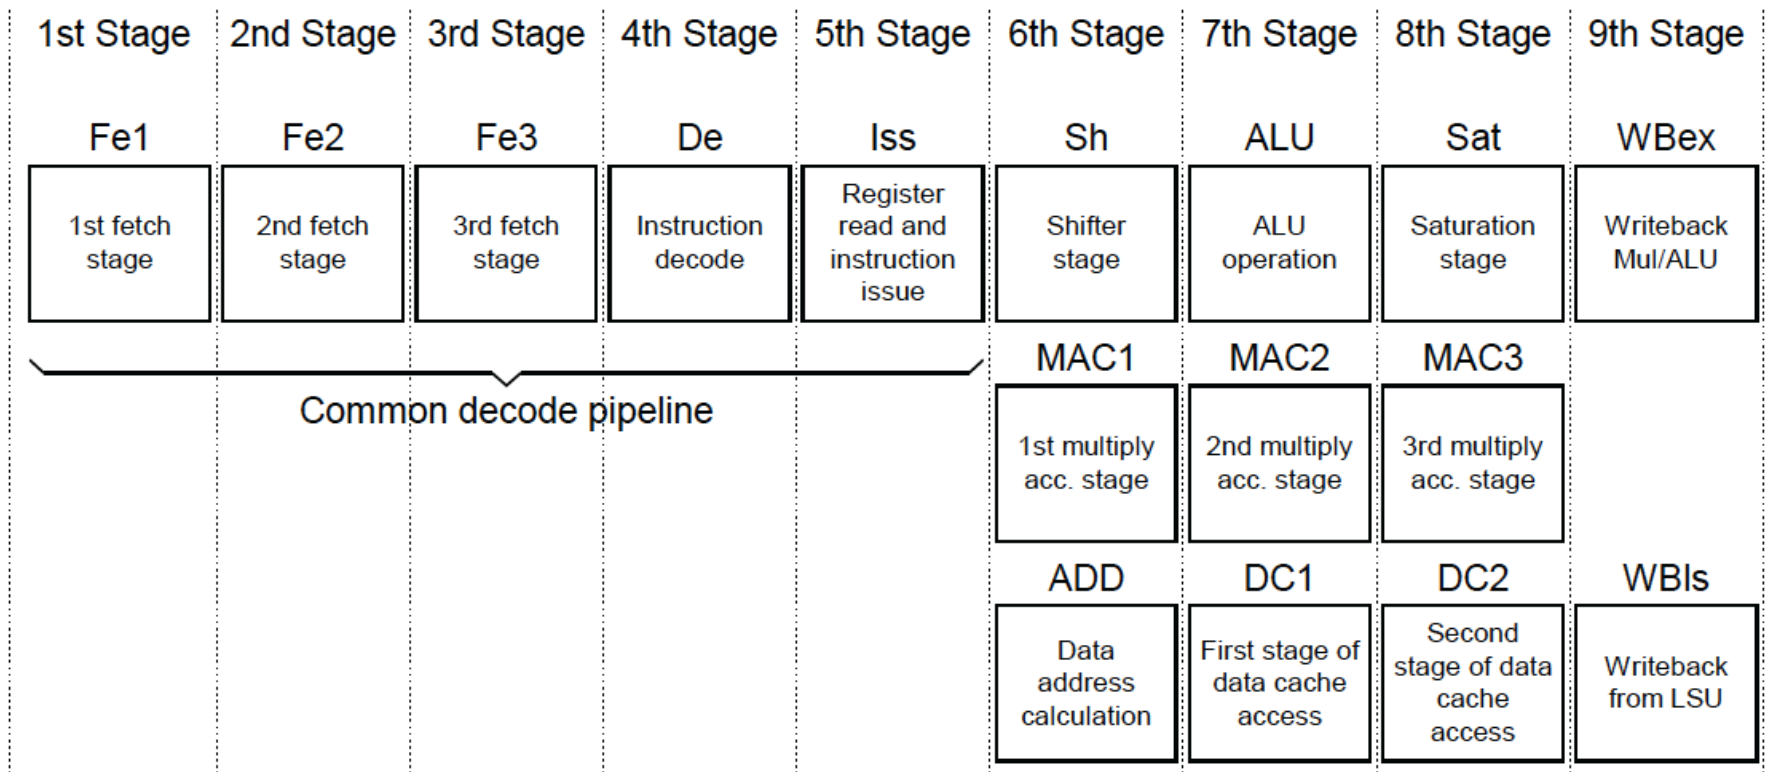
- $TP_{k\text{-stage}} = 1 / (T/k + S)$
 - Objective is to **maximize** throughput
- $Cost_{k\text{-stage}} = G + Lk$
 - Objective is to **minimize** cost

$$\frac{Cost}{Perf} = \frac{kL + G}{\frac{1}{\frac{T}{k} + S}} = (kL + G) \left(\frac{T}{k} + S \right) =$$
$$LT + GS + LS k + \frac{GT}{k}$$


$$\frac{d}{dk} \left(\frac{Cost}{Perf} \right) = 0 + 0 + LS - \frac{GT}{k^2}$$

$$LS - \frac{GT}{k^2} = 0 \Rightarrow k_{opt} = \sqrt{\frac{GT}{LS}}$$

ARM Pipeline Stages




Perspective on Deeper Pipeline



Source: Ed Grochowski, 1997

Pitfall

- $TP_{k\text{-stage}} = 1 / (T/k + S)$
 - We have assumed that the stages are all balanced
 - What if they are not?
 - We will see how the throughput deviates from ideal throughput for imbalanced pipeline
 - Throughput computation: $1 / (\text{maximum delay of all stages})$



Pipeline Design Steps

- **Balancing work in pipeline stages**
- **Keeping the pipeline correct, moving, and full in the presence of events that disrupt pipeline flow**
 - Dependencies
 - Resource contentions
 - Long latency operations
- **Handling exceptions, interrupts**
- **Improving pipeline throughput**

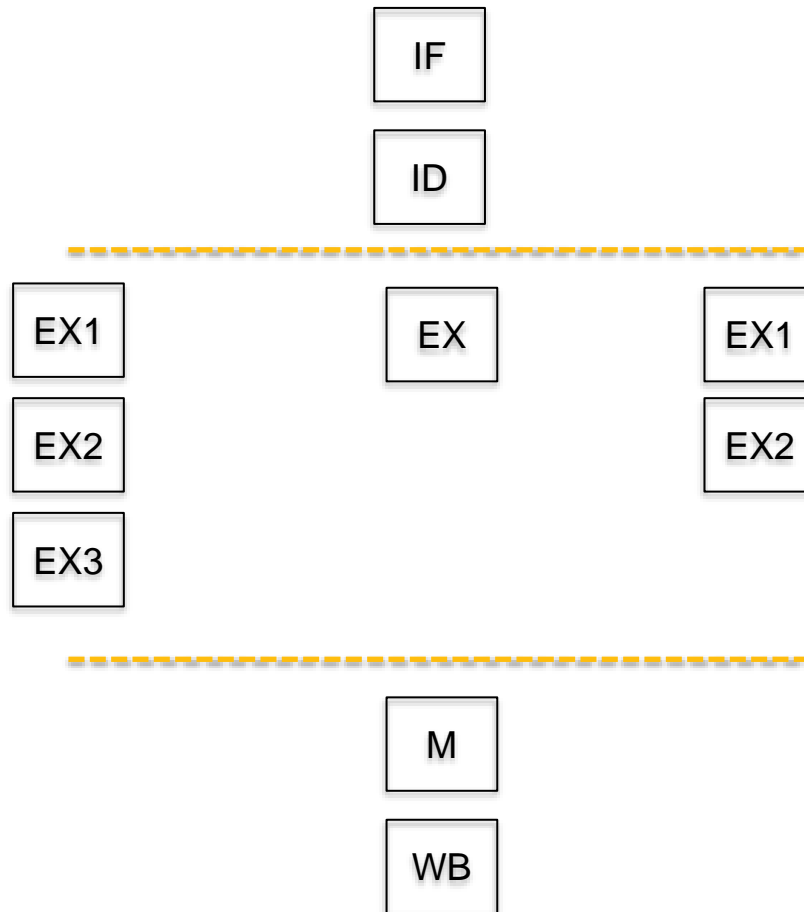
Pipeline Design Steps

- **Balancing work in pipeline stages**
- **Keeping the pipeline correct, moving, and full in the presence of events that disrupt pipeline flow**
 - Dependencies
 - Resource contentions
 - Long latency operations
- **Handling exceptions, interrupts**
- **Improving pipeline throughput**

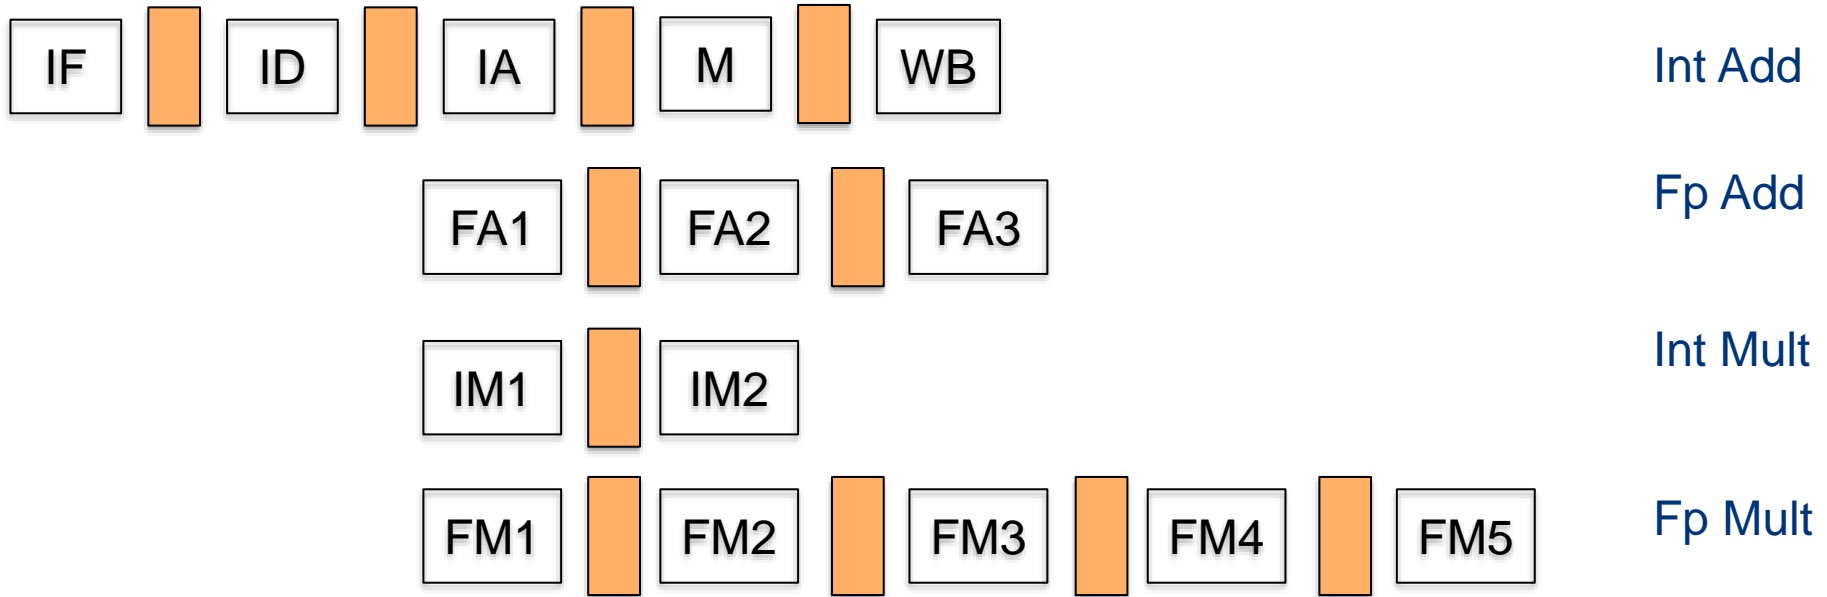
Balancing Pipeline Stages

- **Divide a big operation into small sub-operations**
- **Merge sub-operations**
- **Deeper pipelines**
- **Pipelining of memory accesses**
- **Multiple different pipelines or sub-pipelines**
 - Pipelining of long latency operations

Modern Architecture



Running Pipeline Example



Pipeline Design Steps

- Balancing work in pipeline stages
- Keeping the pipeline correct, moving, and full in the presence of events that disrupt pipeline flow
 - Dependencies
 - Resource contentions
 - Long latency operations
- Handling exceptions, interrupts
- Improving pipeline throughput